

Linux Device Driver

Sunbeam Infotech




IO Ports

GPIO1 → DATAOUT = BV(17); → arm

outl(BV(17), &GPIO1 → DATAOUT); → ldd

STR (arm)
OUT (x86)




Device Driver

HAL - asm code
(Arch Spec) → inb(), outb(), ...
 inb(), outb(), ...

IO: mem mapped
HW: OR IO mapped

e.g. arm
e.g. x86

IO Ports



Memory mapped IO

- IO bus & memory bus is overlapping
- Same instr for mem & IO access
- e.g. ARM, ...
 ↳ LDR, STR.

IO mapped IO

- Different buses/
 Special signal for mem & IO.
- Different instr for mem & IO access.
- e.g. x86, ...
 ↳ IN, OUT



Hardware interaction

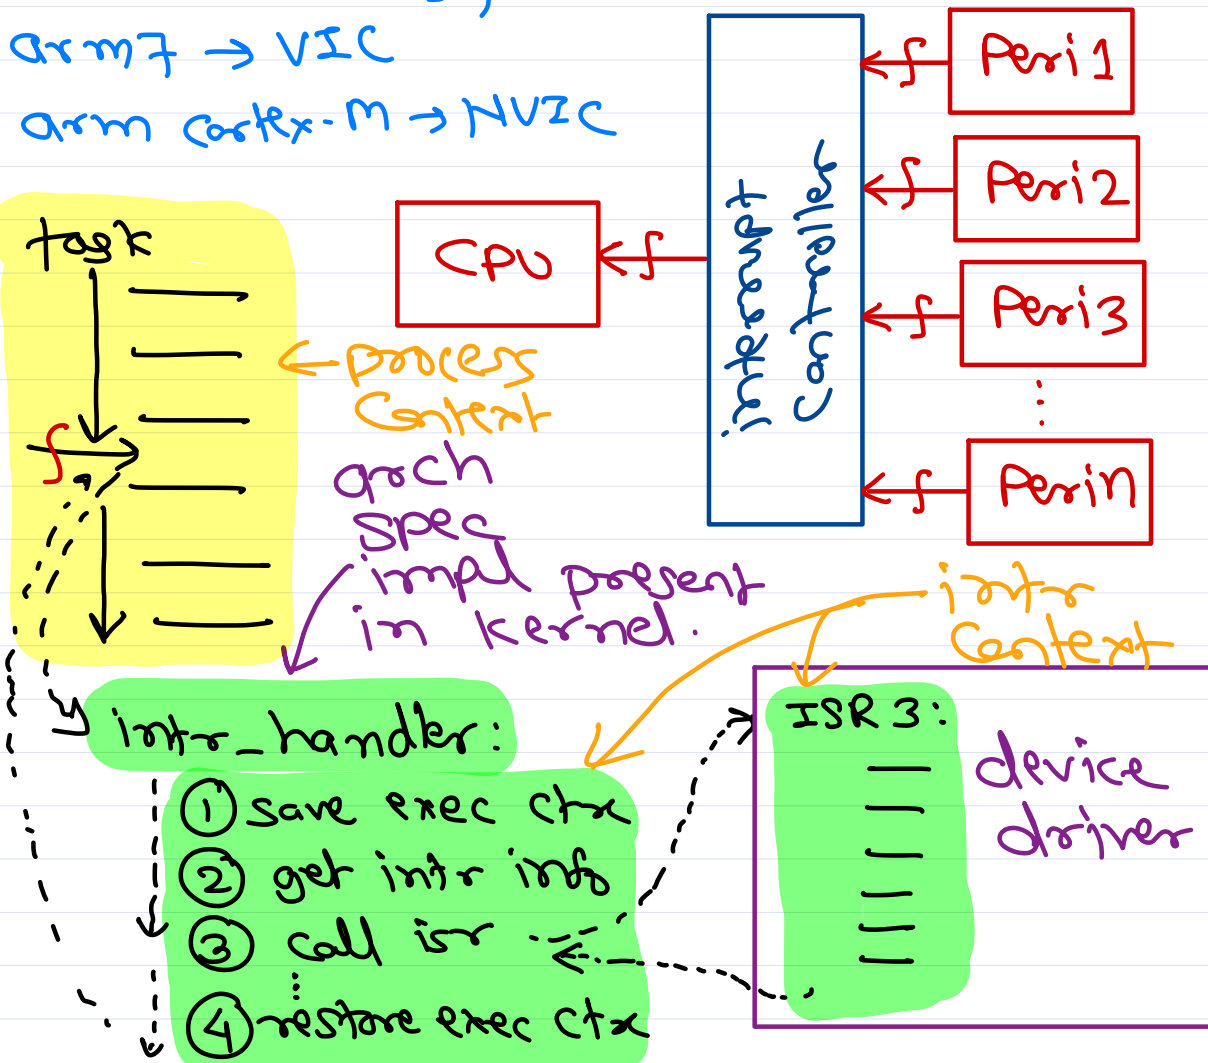
- IO devices are interfaced with CPU via IO ports.
 - On x86 system, this is IO mapped IO.
 - On ARM system, this is memory mapped IO.
- To ensure uniform programming, kernel provides IO access macros/functions in HAL.
- Before accessing IO ~~memory~~^{port} addresses, they should be owned by the driver. This can be done by *request_region()*. It can be released at the end using *release_region()*.
- Actual IO operation can be done using *inb()*, *outb()*, *inw()*, *outw()*, *inl()*, *outl()*, ...
- Device driver should also handle interrupts produced by the hardware device. The ISR is registered using *request_irq()*. It is released using *free_irq()*.
- ISR should not contain blocking code, because ISR runs in interrupt context. Any long running task should be deferred in tasklet, workqueue or timer (as appropriate).
- Typical hardware init and de-init code is done in *open()* and *release()* driver operation; while actual data transfer is done in *read()* and *write()* operation.

cat /proc/ioports



Interrupt Handling in Linux

x86 → PIC-8259
arm7 → VIC
arm cortex-M → NVIC



ISR impl:

irqreturn_t **isr_fn**(int **irq**, void *param):

- IRQ_HANDLED → intr handled by this ISR.
- IRQ_NONE → intr not handled by this ISR.
↳ kernel calls next isr on same line.

Register ISR:

request_irq(**irq**, **isr_fn**, flags, name, param):

cat /proc/interrupts

① IRQF_DISABLED

② IRQF_TIMER

③ IRQF_SAMPLE_RANDOM

④ IRQF_SHARED

↳ true random numbers

↳ /dev/random

↳ random entropy pool

↳ Same irq line is shared for multiple hw devices.

Unregister ISR:

free_irq(**irq**, param):


e.g. com1, com2, com3, com4 ← Serial ports

4 3 4 3



Interrupt

- Interrupts are special signals sent from device to CPU.
- Interrupt handling is architecture specific.



Interrupt handling

- Interrupt is sent from the device to the PIC.
- PIC inform CPU about interrupt through interrupt line.
- CPU pause current task execution and execute interrupt handler.
- Interrupt handler does following
 - Save current task context on stack.
 - Get interrupt details from PIC.
 - Call ISR to handle the interrupt.
 - Invoke scheduler.
 - Restore the task context.
- In Linux there are two execution context.
 - Process context
 - User space process or kernel thread context. May block.
 - Interrupt context
 - Interrupt handler and ISR execution context.
 - Atomic context: cannot block.



Interrupt handling in Linux

- Since interrupt context cannot block, handler/ISR should return immediately.
- Heavy processing and/or blocking tasks should be deferred.
- Linux divides interrupt handling into two parts
 - Top half → *ISR → Interrupt Context*
 - Run immediately when interrupt arrives.
 - Do time critical and non-blocking task like interrupt acknowledgement.
 - Cannot be pre-empted by another interrupt from same device.
 - Bottom half → *Soft IRQ, Tasklet, Work Queue*
 - Variety of bottom half implementations in Linux kernel.
 - Execute later – in interrupt context or process context.
 - Do heavy processing and/or blocking tasks.
 - Can be pre-empted by interrupt (top-half).
- Interrupt handling must be done in corresponding device driver.
 - Driver should implement top-half and/or bottom-half as per requirement.
 - Linux kernel ensure uniform programming model irrespective of architecture.



Implementing top half

- Two step process

- Implement ISR.
- Register ISR.

- ISR registration

- `#include <linux/interrupt.h>`

- `int request_irq(unsigned int irq, irq_handler_t handler, unsigned long flags, const char *name, void *dev);`

- irq: interrupt number
- handler: typedef `irqreturn_t (*irq_handler_t)(int, void *)`;
- flags:

- `IRQF_DISABLED`
- `IRQF_SAMPLE_RANDOM`
- `IRQF_TIMER`
- `IRQF_SHARED`

- name: device name /proc/irq and /proc/interrupts
- dev: extra information to be passed to handler.
- Returns 0 on success or `-EBUSY` if interrupt line is already in use.

- `request_irq()` may block and should not be called from interrupt context. Typically called when opening the device for processing or module initialization.

`irqreturn_t my_isr(int irq, void *param) {`
`request_irq(`
`3`

multiple device instances - private obj to keep each device info
e.g. struct serial_info { struct private_struct devices[4];
int irq;
int ioaddr;
mutex m;
...

`com1`
`request_irq(4, my_isr, IRQF_SHARED, "com1", &devices[0]);`
`3;`
`com3`
`request_irq(4, my_isr, IRQF_SHARED, "com3", &devices[2]);`

irq=4 io=0x3F8 : com1	0
irq=3 io=0x2F8 : com2	1
irq=4 io=0x3E8 : com3	2
irq=3 io=0x2E8 : com4	3




Implementing top half

- ISR un-registration
 - Interrupt line must be released while unloading module or closing device.
 - void free_irq(unsigned int irq, void *dev);
- Implementing ISR
 - irqreturn_t my_intr_handler(int irq, void *dev);
 - irq: interrupt number
 - dev: extra param passed while request_irq()
 - returns IRQ_HANDLED or IRQ_NONE.
 - Should contain time-critical tasks and interrupt acknowledgement.
 - Also trigger bottom-half if required.
 - Should not sleep/block.
- Linux interrupt handlers are not re-entrant. Current interrupt line is disabled while execution of ISR.
- Shared interrupt handlers
 - Must pass unique dev param – typically device private struct.
 - ISR must check if interrupt is raised from the corresponding device before handling it.
 - Kernel execute all ISR registered on same interrupt line.



Interrupt handling

- Interrupt context
 - Atomic context.
 - One page kernel stack per processor.
- Interrupt execution




Interrupt control

- ✗ local_irq_disable(): Disables local interrupt delivery
 - ✗ local_irq_enable(): Enables local interrupt delivery
 - ✓ local_irq_save(): Saves the current state of local interrupt delivery and then disables it
 - ✓ local_irq_restore(): Restores local interrupt delivery to the given state
 - disable_irq(): Disables the given interrupt line and ensures no handler on the line is executing
 - enable_irq(): Enables the given interrupt line
 - irqs_disabled(): Returns nonzero if local interrupt delivery is disabled; otherwise returns zero
 - in_interrupt(): Returns nonzero if in interrupt context and zero if in process context
 - in_irq(): Returns nonzero if currently executing an interrupt handler and zero otherwise
- ISR

← Current CPU



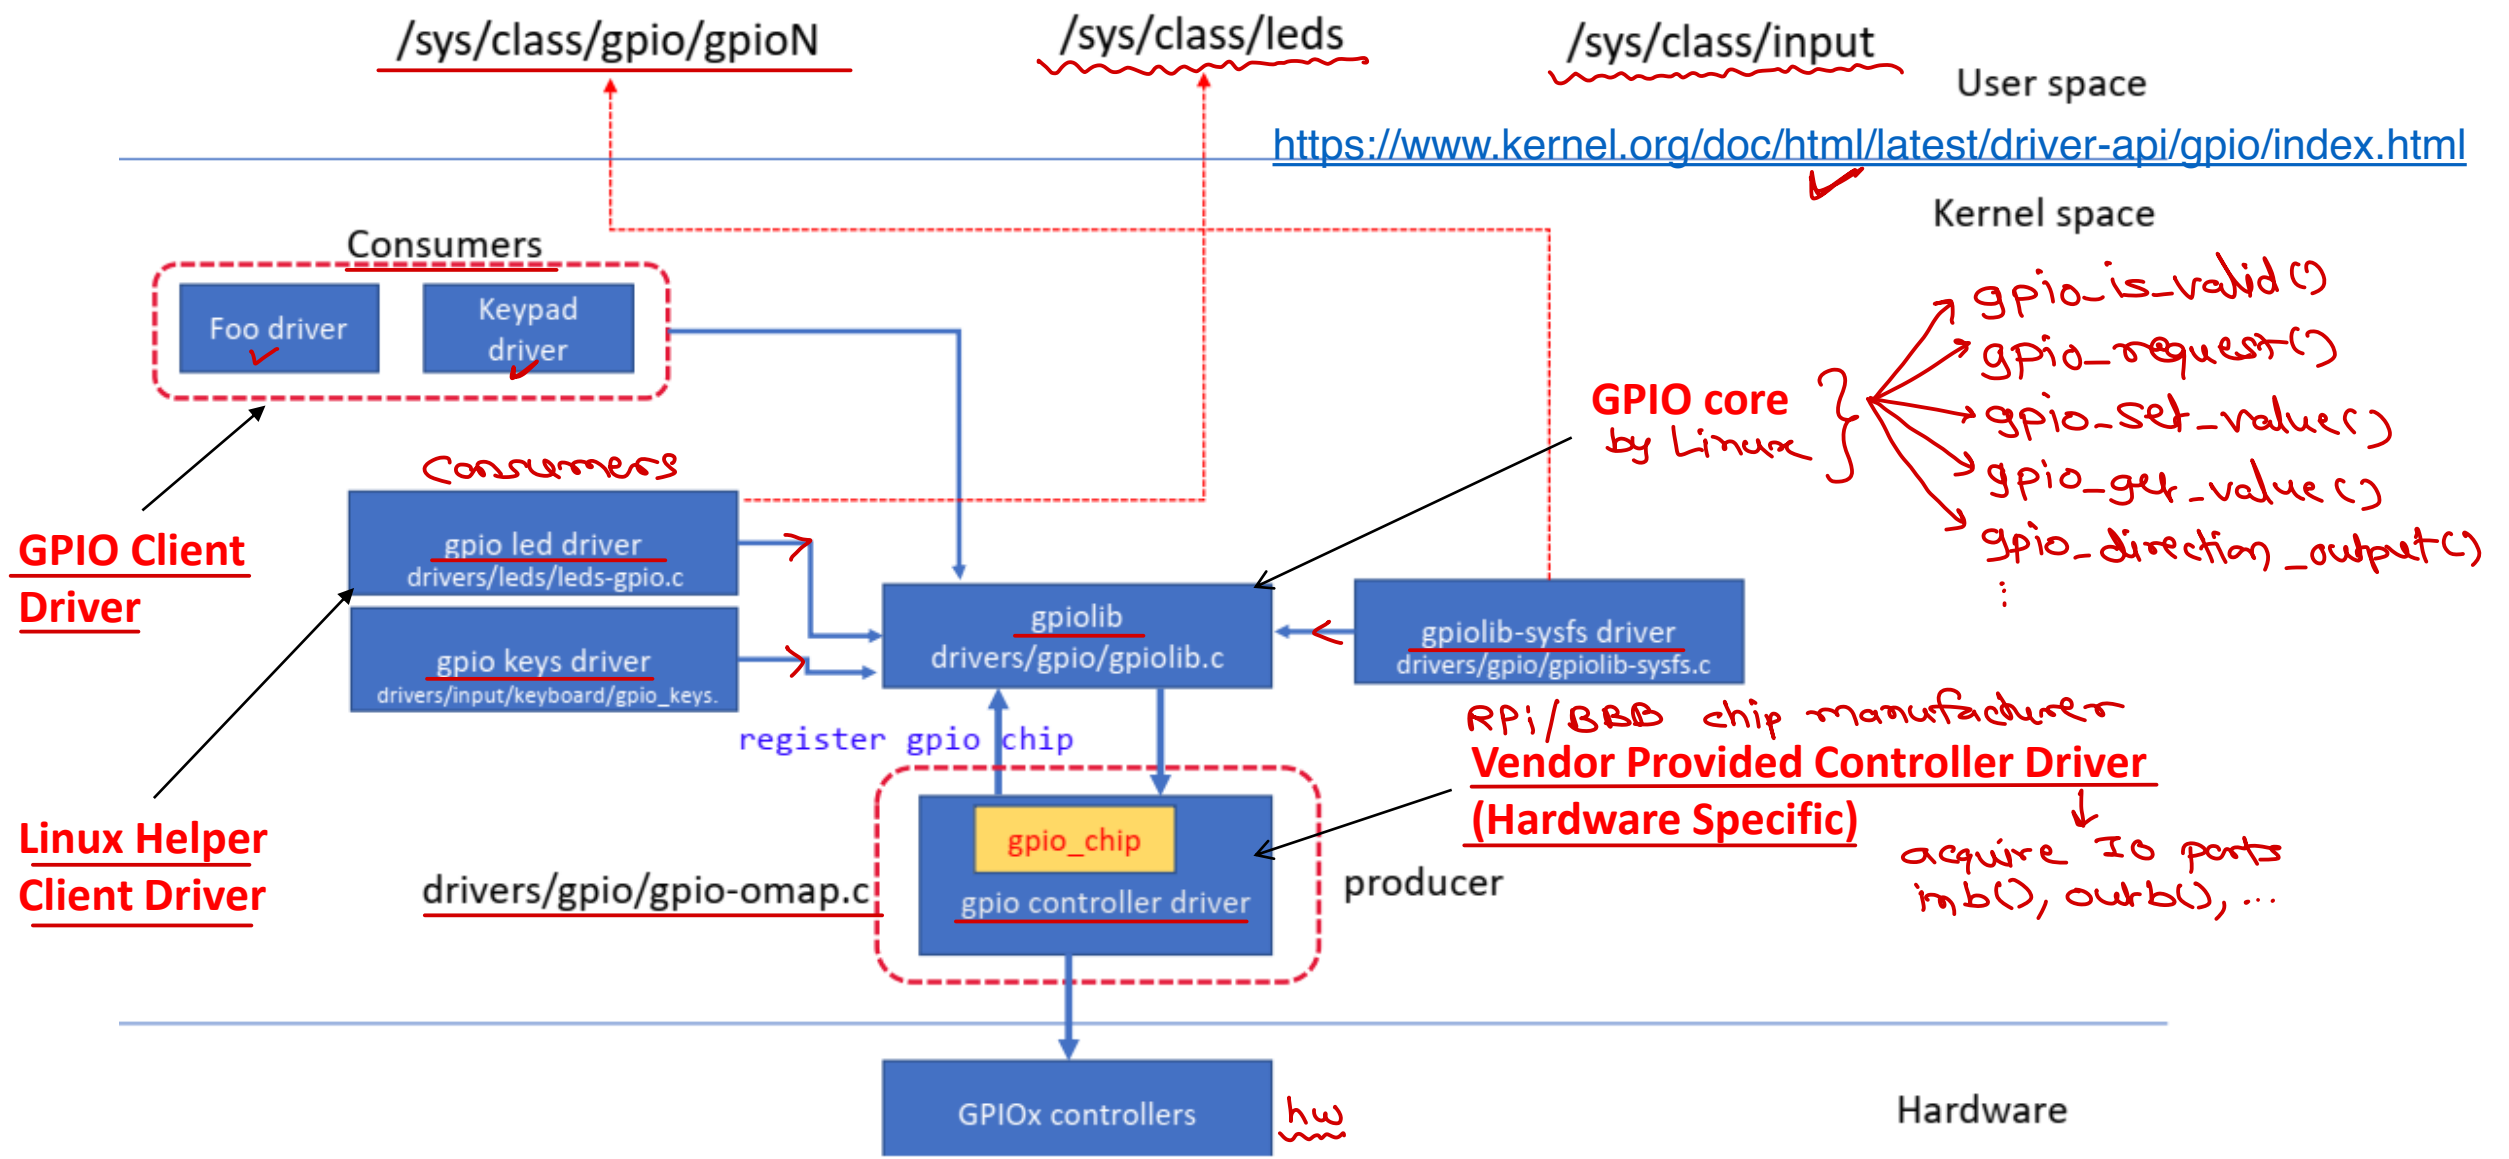


Linux GPIO SubSystem

Sunbeam Infotech



Linux GPIO SubSystem - producer/consumer pattern



Using Linux GPIO subsystem

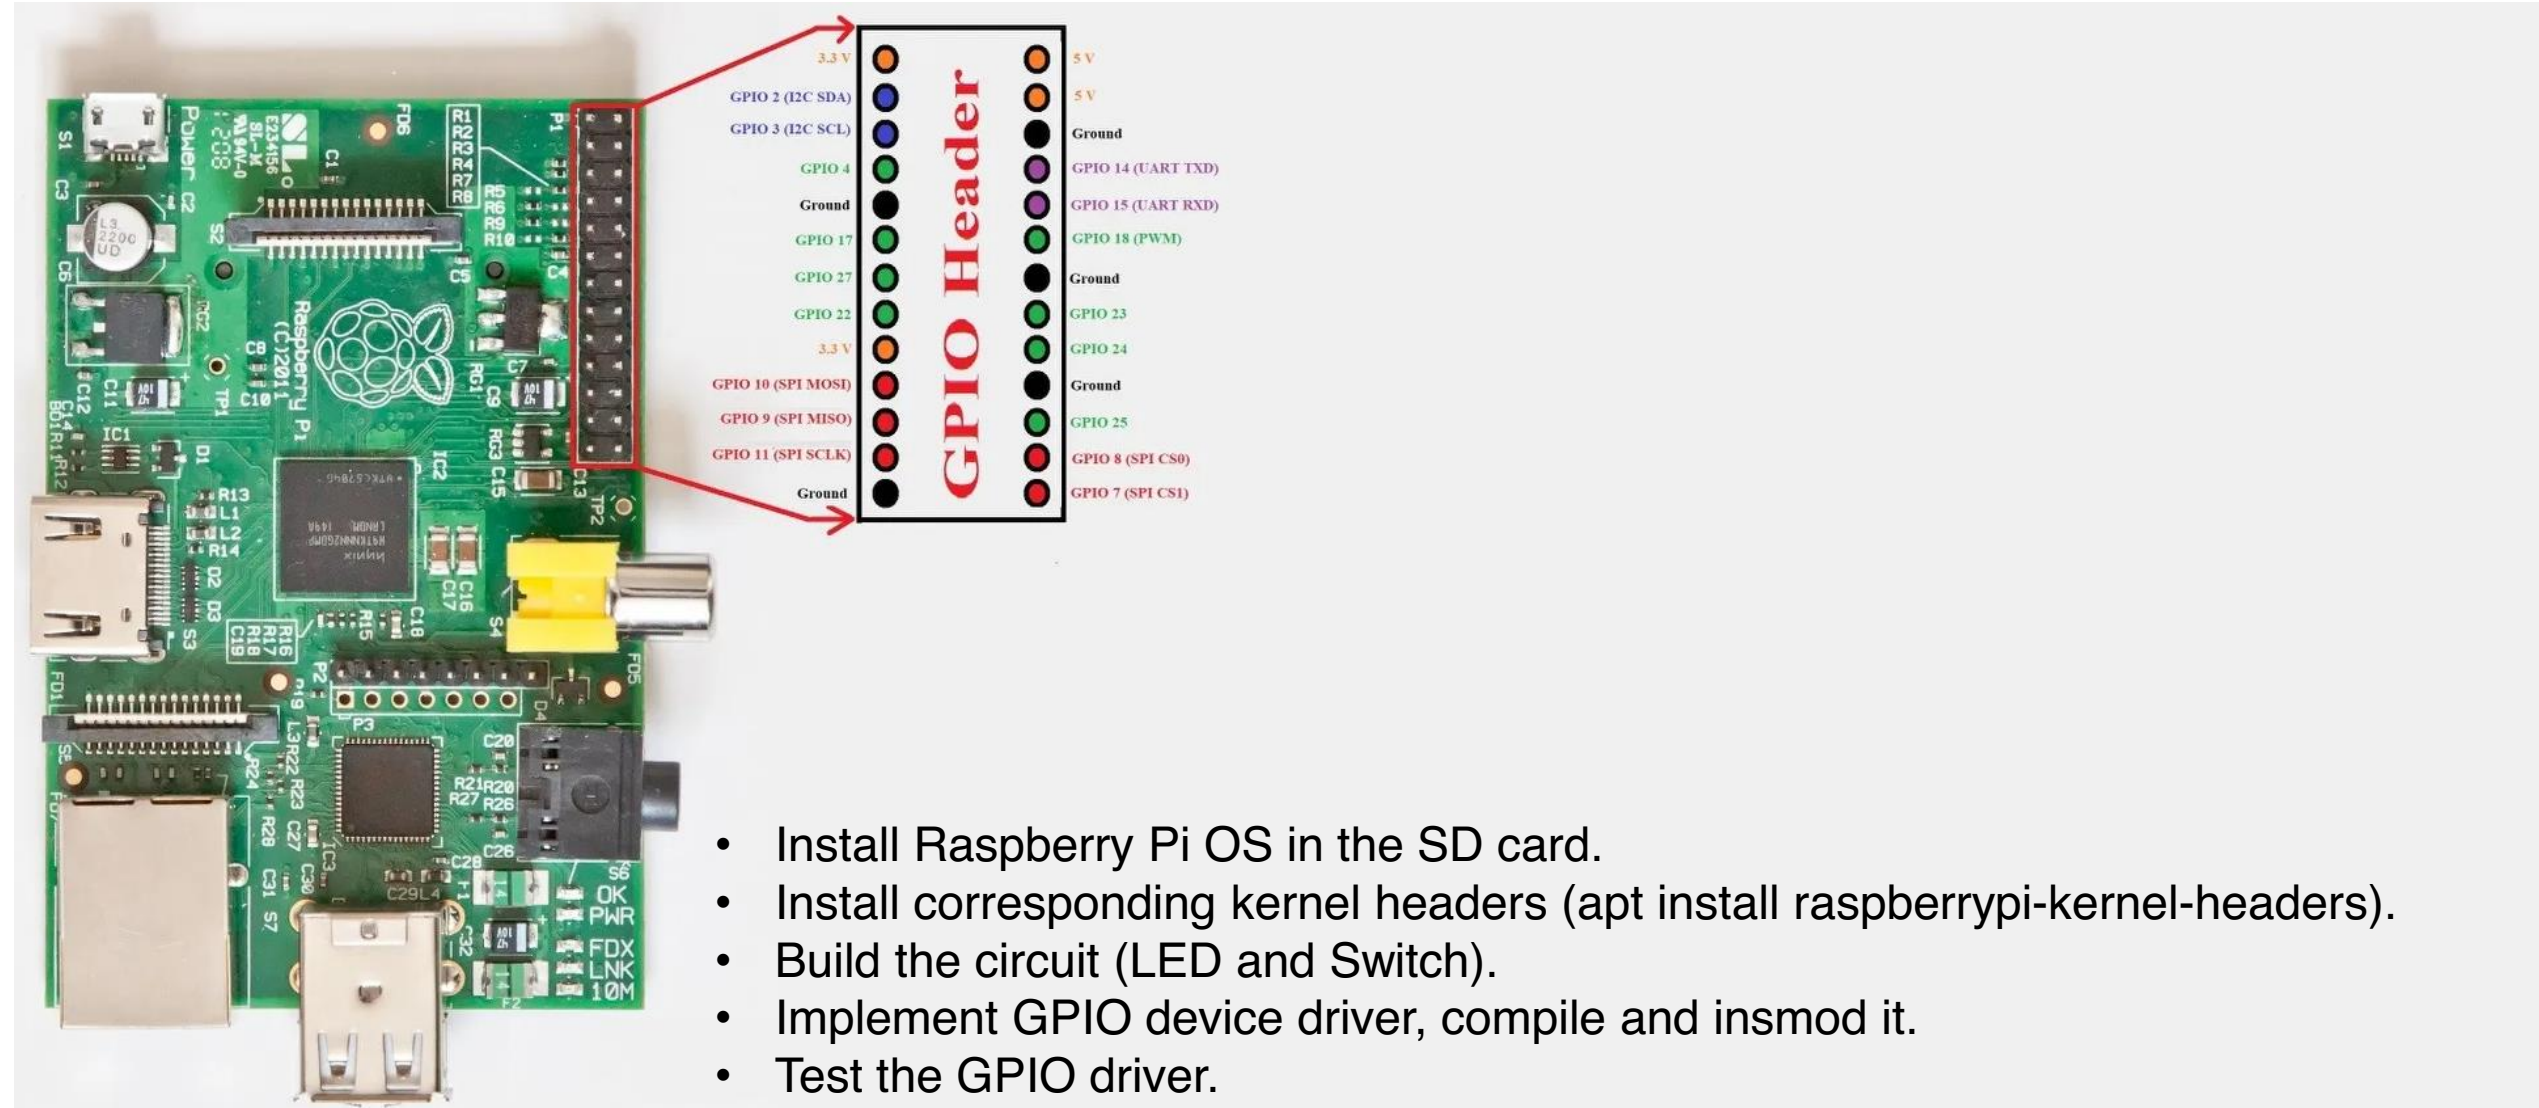
- Verify the GPIO is valid or not. `bool gpio_is_valid(int gpio_number);`
- If valid, request the GPIO from the Kernel GPIO subsystem. `int gpio_request(unsigned gpio, const char *label);`
 - `int gpio_request_one(unsigned gpio, unsigned long flags, const char *label);` – Request one GPIO.
 - `int gpio_request_array(struct gpio *array, size_t num);` – Request multiple GPIOs.
- Export GPIO to sysfs. `int gpio_export(unsigned int gpio, bool direction_may_change); void gpio_unexport(unsigned int gpio);`
- Set the direction of the GPIO (IN/OUT). optional
 - `int gpio_direction_input(unsigned gpio);` → switch
 - `int gpio_direction_output(unsigned gpio, int initial_value);` → led
- Make the GPIO to High/Low if it is set as an output pin.
 - `gpio_set_value(unsigned int gpio, int value);`
- Set the debounce-interval and read the state if it is set as an input pin. Enable IRQ for edge/level triggered.
 - `int gpio_get_value(unsigned gpio);`
 - `int gpiod_set_debounce(unsigned gpio, unsigned debounce);`
 - `int gpio_to_irq(unsigned gpio);`
 - `request_irq()` with flag `IRQF_TRIGGER_RISING`, `IRQF_TRIGGER_FALLING`, `IRQF_TRIGGER_HIGH`, or `IRQF_TRIGGER_LOW` and `free_irq()`;
- Release the GPIO while exiting the driver. `void gpio_free(unsigned int gpio);`
 - `void gpio_free_array(struct gpio *array, size_t num);` – Release multiple GPIOs.



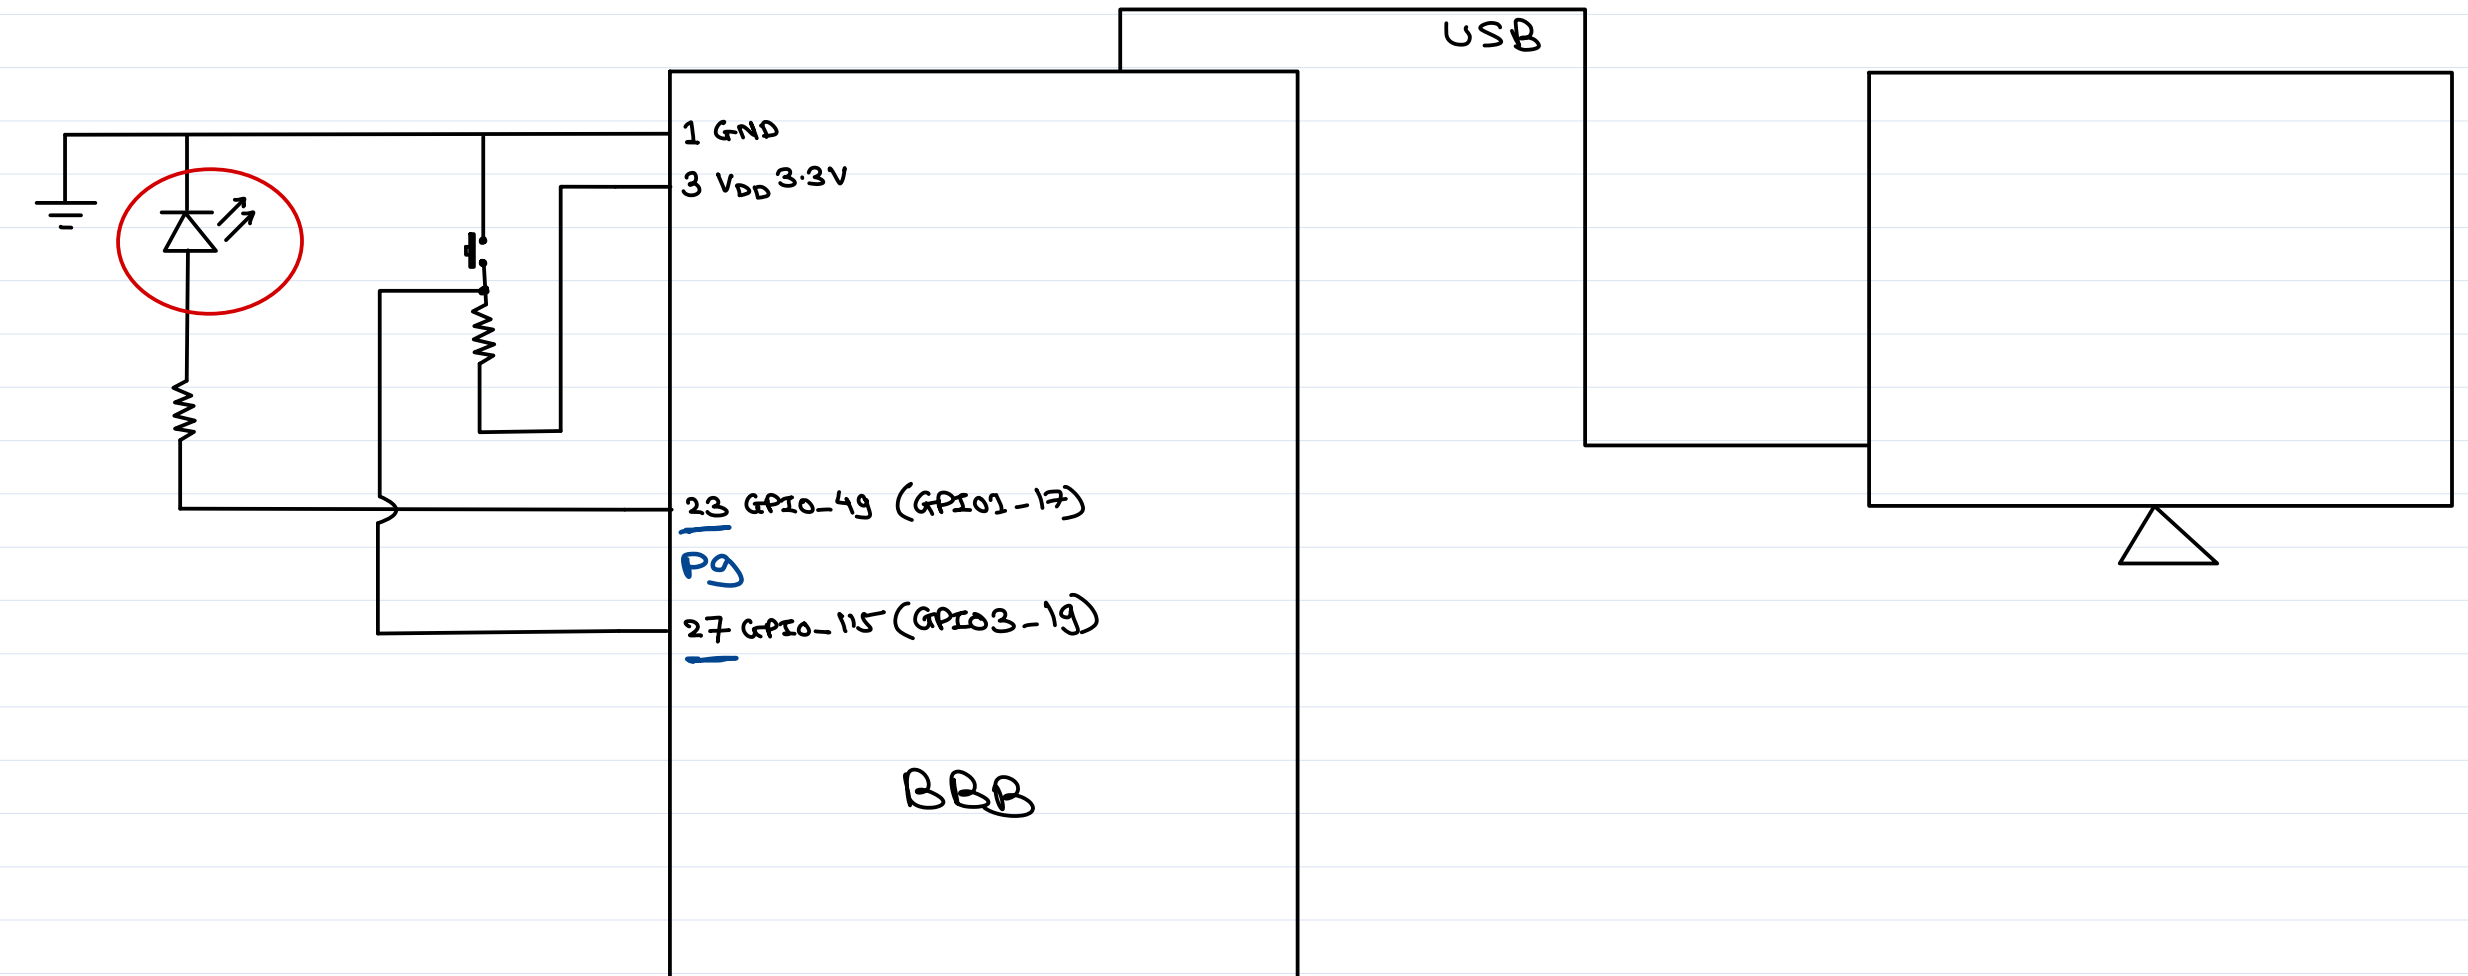
GPIO device driver example



GPIO driver on RPi-1



BBB Led & Switch





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>

