


Linux Device Driver

Sunbeam Infotech



Interrupt Handling in Linux

x86 → PIC-8259
arm7 → VIC
arm cortex-M → NVIC



ISR impl:

irqreturn_t **isr_fn**(int **irq**, void *param):

- IRQ_HANDLED → intr handled by this ISR.
- IRQ_NONE → intr not handled by this ISR.
↳ kernel calls next isr on same line.

Register ISR:

request_irq(**irq**, **isr_fn**, flags, name, param):

cat /proc/interrupts

- ① IRQF_DISABLED
- ② IRQF_TIMER
- ③ IRQF_SAMPLE_RANDOM
- ④ IRQF_SHARED

↳ true random numbers
↳ /dev/random

↳ random entropy pool

↳ Same irq line is shared for multiple hw devices.

Unregister ISR:

free_irq(**irq**, param):

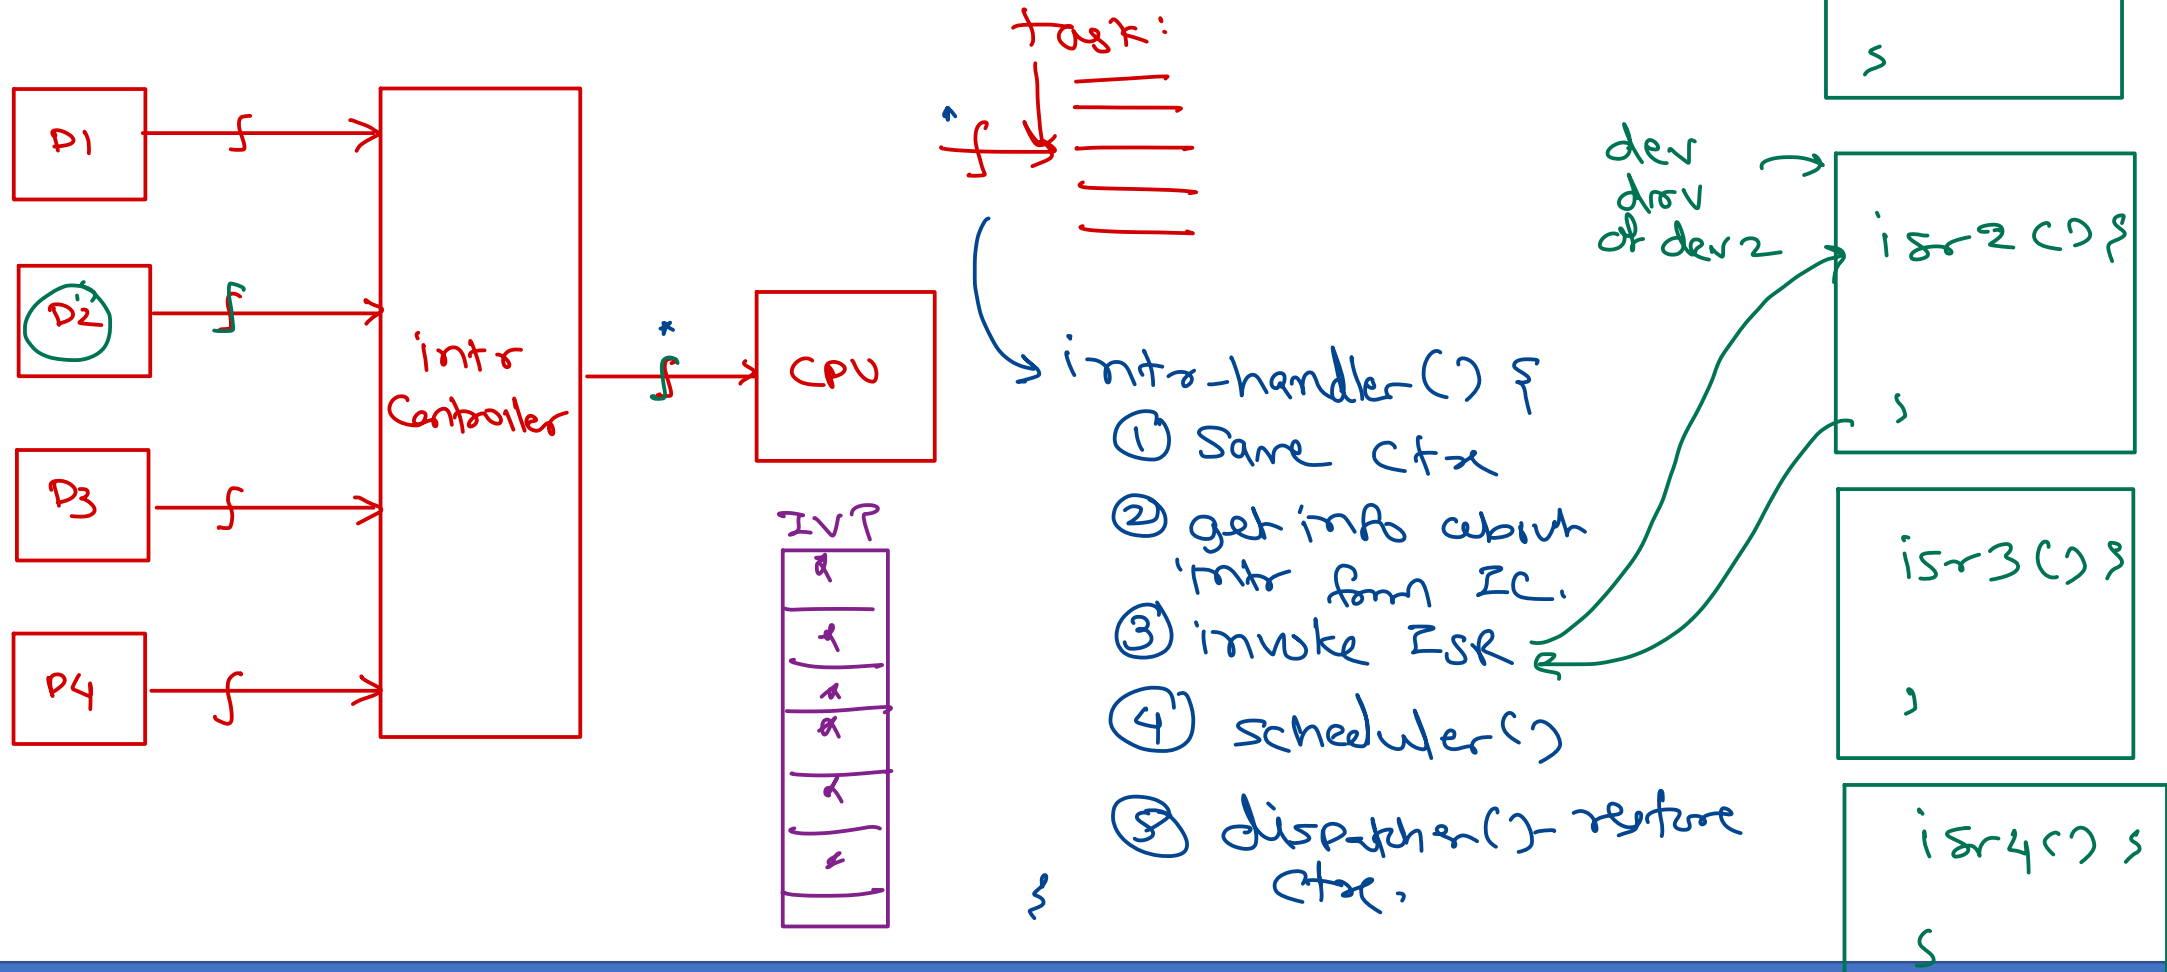
e.g. com1, com2, com3, com4 ← Serial ports

4 3 4 3



Interrupt

- Interrupts are special signals sent from device to CPU.
- Interrupt handling is architecture specific.



Interrupt handling

- Interrupt is sent from the device to the PIC.
- PIC inform CPU about interrupt through interrupt line.
- CPU pause current task execution and execute interrupt handler.
- Interrupt handler does following
 - Save current task context on stack.
 - Get interrupt details from PIC.
 - Call ISR to handle the interrupt.
 - Invoke scheduler.
 - Restore the task context.
- In Linux there are two execution context.
 - Process context
 - User space process or kernel thread context. May block.
 - Interrupt context
 - Interrupt handler and ISR execution context.
 - Atomic context: cannot block.



Interrupt handling in Linux

- Since interrupt context cannot block, handler/ISR should return immediately.
- Heavy processing and/or blocking tasks should be deferred.
- Linux divides interrupt handling into two parts
 - Top half → *ISR → Interrupt Context*
 - Run immediately when interrupt arrives.
 - Do time critical and non-blocking task like interrupt acknowledgement.
 - Cannot be pre-empted by another interrupt from same device.
 - Bottom half → *Soft IRQ, Tasklet, Work Queue*
 - Variety of bottom half implementations in Linux kernel.
 - Execute later – in interrupt context or process context.
 - Do heavy processing and/or blocking tasks.
 - Can be pre-empted by interrupt (top-half).
- Interrupt handling must be done in corresponding device driver.
 - Driver should implement top-half and/or bottom-half as per requirement.
 - Linux kernel ensure uniform programming model irrespective of architecture.



Implementing top half

- Two step process

- Implement ISR.
- Register ISR.

- ISR registration

- `#include <linux/interrupt.h>`

- `int request_irq(unsigned int irq, irq_handler_t handler, unsigned long flags, const char *name, void *dev);`

- irq: interrupt number
- handler: typedef `irqreturn_t (*irq_handler_t)(int, void *)`;
- flags:
 - `IRQF_DISABLED`
 - `IRQF_SAMPLE_RANDOM`
 - `IRQF_TIMER`
 - `IRQF_SHARED`

- name: device name /proc/irq and /proc/interrupts
- dev: extra information to be passed to handler.
- Returns 0 on success or `-EBUSY` if interrupt line is already in use.

- `request_irq()` may block and should not be called from interrupt context. Typically called when opening the device for processing or module initialization.

`irqreturn_t my_isr(int irq, void *param) {`
`request_irq(`
`3`

multiple device instances - private obj to keep each device info
e.g. struct serial_info { struct private_struct devices[4];
int irq;
int ioaddr;
mutex m;
...
};

`com1`
`request_irq(4, my_isr, IRQF_SHARED, "com1", &devices[0]);`
`com2`
`request_irq(3, my_isr, IRQF_SHARED, "com2", &devices[1]);`
`com3`
`request_irq(4, my_isr, IRQF_SHARED, "com3", &devices[2]);`
`com4`

irq=4	io=0x3F8	0
:	com1	
irq=3	io=0x2F8	1
:	com2	
irq=4	io=0x3E8	2
:	com3	
irq=3	io=0x2E8	3
:	com4	



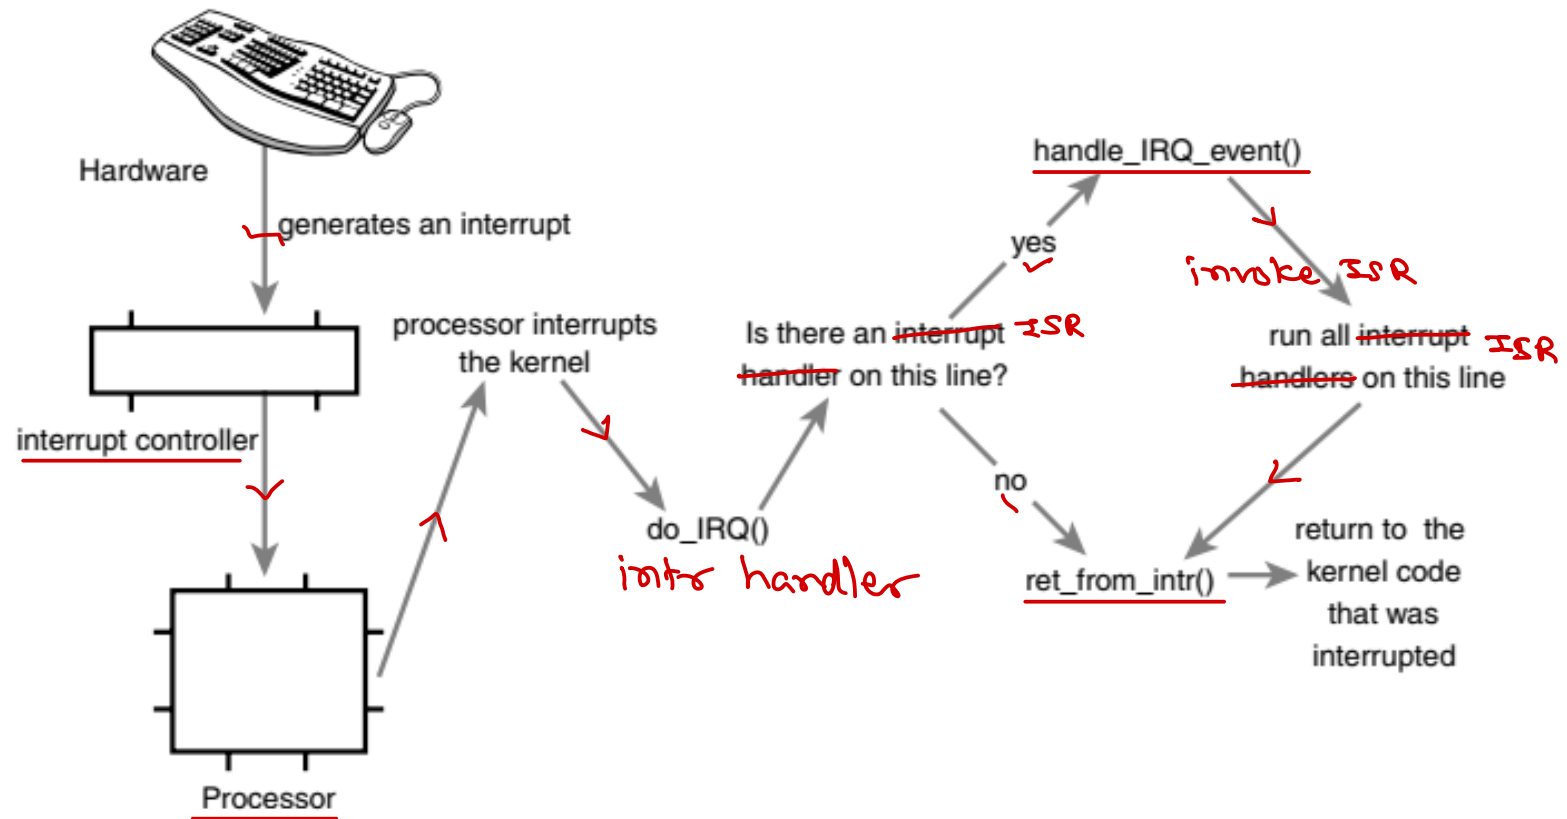
Implementing top half

- ISR un-registration
 - Interrupt line must be released while unloading module or closing device.
 - void free_irq(unsigned int irq, void *dev);
- Implementing ISR
 - irqreturn_t my_intr_handler(int irq, void *dev);
 - irq: interrupt number
 - dev: extra param passed while request_irq()
 - returns IRQ_HANDLED or IRQ_NONE.
 - Should contain time-critical tasks and interrupt acknowledgement.
 - Also trigger bottom-half if required.
 - Should not sleep/block.
- Linux interrupt handlers are not re-entrant. Current interrupt line is disabled while execution of ISR.
- Shared interrupt handlers
 - Must pass unique dev param – typically device private struct.
 - ISR must check if interrupt is raised from the corresponding device before handling it.
 - Kernel execute all ISR registered on same interrupt line.



Interrupt handling

- Interrupt context
 - Atomic context.
 - One page kernel stack per processor.
- Interrupt execution



Interrupt control

- ✗ local_irq_disable(): Disables local interrupt delivery
 - ✗ local_irq_enable(): Enables local interrupt delivery
 - ✓ local_irq_save(): Saves the current state of local interrupt delivery and then disables it
 - ✓ local_irq_restore(): Restores local interrupt delivery to the given state
 - disable_irq(): Disables the given interrupt line and ensures no handler on the line is executing
 - enable_irq(): Enables the given interrupt line
 - irqs_disabled(): Returns nonzero if local interrupt delivery is disabled; otherwise returns zero
 - in_interrupt(): Returns nonzero if in interrupt context and zero if in process context
 - in_irq(): Returns nonzero if currently executing an interrupt handler and zero otherwise
- ISR

← Current CPU



Bottom half

- If interrupt handling need to do heavy processing or blocking task, then driver must implement it in bottom half.
- General guideline for top and bottom half work division:
 - top half {
 - If the work is time sensitive, perform it in the interrupt handler. *ISR*
 - If the work is related to the hardware, perform it in the interrupt handler. *ISR*
 - If the work needs to ensure that another interrupt doesn't interrupt it, perform it in the interrupt handler. *ISR*
 - For everything else, consider performing the work in the bottom half.
- Bottom half are executed after top half.
 - Immediately after top half in interrupt context.
 - In some specialized process context, when no other another high priority task is running.
- Types of bottom halves
 - BH - Removed in kernel 2.5
 - Task queue - Removed in kernel 2.5
 - Softirq - Added in kernel 2.3
 - Tasklet - Added in kernel 2.3
 - Work queue - Added in kernel 2.3



Softirq

- Softirqs are statically allocated at compile time.

```
struct softirq_action {  
    void (*action)(struct softirq_action *);  
};  
static struct softirq_action softirq_vec[NR_SOFTIRQS];
```

↗ 32

- Softirqs are implemented for specialized sub-systems.
- Kernel 2.6.34 have 9 Softirqs implemented.

•	HI_SOFTIRQ	0	High-priority tasklets
•	TIMER_SOFTIRQ	1	Timers
•	NET_TX_SOFTIRQ	2	Send network packets
•	NET_RX_SOFTIRQ	3	Receive network packets
•	BLOCK_SOFTIRQ	4	Block devices
•	TASKLET_SOFTIRQ	5	Normal priority tasklets
•	SCHED_SOFTIRQ	6	Scheduler
•	HRTIMER_SOFTIRQ	7	High-resolution timers
•	RCU_SOFTIRQ	8	RCU locking



Softirqs

- Softirqs must be triggered for the execution. This is called as “raising Softirq”. Mostly done from ISR.
- Pending Softirq are checked and executed in one of the following place – do_softirq().
 - In the return from hardware interrupt code path
 - In the ksoftirqd kernel thread (per processor)
 - In any code that explicitly checks for and executes pending softirqs, such as the networking subsystem
- Using Softirq
 - Currently only network and block subsystem is using Softirq directly.
 - It is not advised to use Softirqs directly.
 - Softirq can be registered using open_softirq() and can be triggered using raise_softirq().

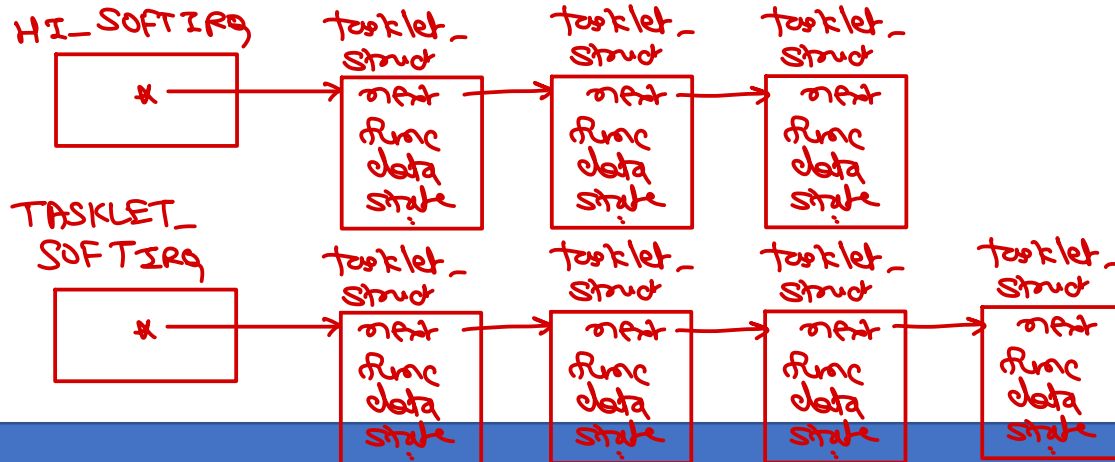


Tasklets

- Implemented on top of Softirqs i.e. HI_SOFTIRQ^① and TASKLET_SOFTIRQ^⑤.
- Tasklets are dynamic components and much easier to use.

```
struct tasklet_struct {  
    struct tasklet_struct *next; /* next tasklet in the list */  
    unsigned long state; /* state of the tasklet */  
    atomic_t count; /* reference counter */  
    void (*func)(unsigned long); /* tasklet handler function */  
    unsigned long data; /* argument to the tasklet function */  
};
```

- Tasklet state can be: 0, TASKLET_STATE_SCHED or TASKLET_STATE_RUN.



Tasklets

- Declare Tasklet statically.
 - DECLARE_TASKLET(my_tasklet, my_tasklet_handler, dev);
- Declare & initialize Tasklet dynamically.
 - struct tasklet_struct my_tasklet;
 - tasklet_init(t, tasklet_handler, dev); *arg to tasklet func.*
- Tasklet handler implementation
 - void tasklet_handler(unsigned long data) { ... }
 - Like softirq, tasklet cannot sleep/block.
 - While executing tasklet, interrupts are enabled.
 - Tasklet are not re-entrant or execute concurrently.
- Trigger Tasklet
 - tasklet_schedule(&my_tasklet); *tasklet_hi_schedule(&my_tasklet);*
 - Change tasklet state to TASKLET_STATE_SCHED.
- Tasklet can be enabled/disabled explicitly.
 - tasklet_enable(&my_tasklet);
 - tasklet_disable(&my_tasklet);



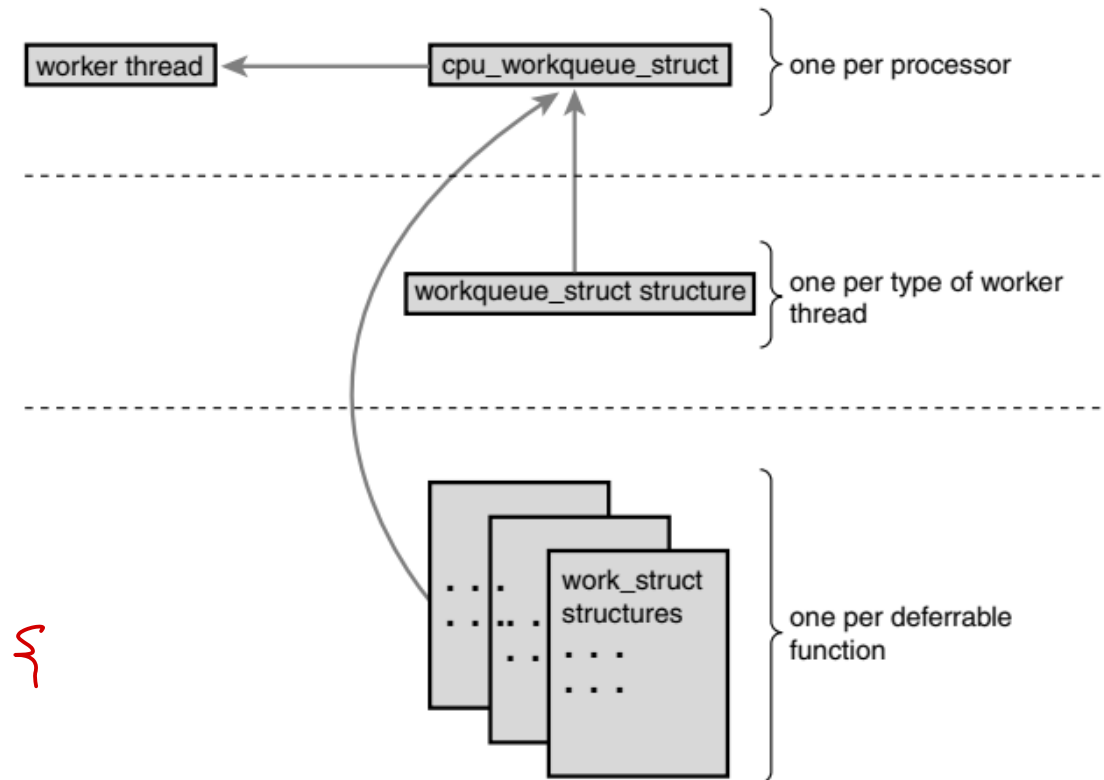
Work queue

- Work queues defer work into a kernel thread. *→ kworker*
- Always runs in process context – worker threads (per processor).
- Work queues are schedulable and can sleep/block.
- Usual alternative to work queues is kernel threads. However creating new kernel threads isn't advised.

```
struct cpu_workqueue_struct {  
    spinlock_t lock; /* lock protecting this structure */  
    struct list_head worklist; /* list of work */  
    wait_queue_head_t more_work;  
    struct work_struct *current_struct;  
    struct workqueue_struct *wq; /* associated */  
    task_t *thread; /* associated thread */  
};
```

```
struct work_struct {  
    atomic_long_t data;  
    struct list_head entry;  
    work_func_t func;  
};
```

void func (data) {



Work queue

- Creating work
 - DECLARE_WORK(name, void (*func)(void *), void *data); // static
 - INIT_WORK(struct work_struct *work, void (*func)(void *), void *data); // dynamic
- Work handler
 - void work_handler(void *data) { ... } → may sleep
- Scheduling work
 - schedule_work(&work);
 - schedule_delayed_work(&work, delay);
- Ensure work completion
 - void flush_scheduled_work(void);
- Cancel scheduled work
 - int cancel_delayed_work(struct work_struct *work);




Control bottom half

- To enable/disable all bottom half (tasklet & work queue) processing
 - `void local_bh_enable();`
 - `void local_bh_disable();`
- Choosing bottom half
 - Softirq
 - Concurrent execution and need synchronization
 - Good for fast execution and high frequency use
 - Cannot sleep
 - Tasklet
 - Simplified programming
 - Not executed concurrently
 - Cannot sleep
 - Work queue
 - Can sleep
 - Higher overheads due to kernel thread & context switching



Kernel memory management

RAM (physical memory)




x86-32 architecture

Zone wise buddy allocator is employed to allocate contiguous physical pages as per requirement.
Buddy allocator keeps track of free pages.




Buddy allocator



① 128 → C

- 1024 →
- 512 → A'
- 256 → B'
- 128 → C'
- 64 →
- 32 →


Buddy allocator



- ① 128 → **C**
- ② 64 → **D**

- 1024 →
- 512 → **A'**
- 256 → **B'**
- 128 →
- 64 → **D'**
- 32 →


Buddy allocator



- ① 128 → C
- ② 64 → D
- ③ 128 → E

- 1024 →
- 512 → A'
- 256 →
- 128 → E'
- 64 → D'
- 32 →


Buddy allocator




- ① 128 → C
- ② 64 → D
- ③ 128 → E
- ④ 256 → F

- 1024 →
- 512 →
- 256 → F'
- 128 → E'
- 64 → D'
- 32 →

Buddy allocator



Buddy allocator



Allocations


- ① 128 → C
- ② 64 → D
- ③ 128 → E
- ④ 256 → F
- ⑤ 32 → G

Deallocation

- ① 64 - D

- 1024 →
- 512 →
- 256 → F'
- 128 → E'
- 64 → D
- 32 → G'

Buddy allocator



Allocations


- ① 128 → C
- ② 64 → D
- ③ 128 → E
- ④ 256 → F
- ⑤ 32 → G

Deallocation

- ① 64 - D
- ② 128 - C

- 1024 →
- 512 →
- 256 → F'
- 128 → E', C
- 64 → D
- 32 → G'

Buddy allocator




Allocations


- ① 128 → C
- ② 64 → D
- ③ 128 → E
- ④ 256 → F
- ⑤ 32 → G

Deallocation

- ① 64 - D
- ② 128 - C
- ③ 32 - G



Buddy allocator



Allocations

- ① 128 → C
- ② 64 → D
- ③ 128 → E
- ④ 256 → F
- ⑤ 32 → G
- ⑥ 32 → H

Deallocation

- ① 64 - D
- ② 128 - C
- ③ 32 - G

For this allocation, we need to again break higher order blocks, which were merged in earlier request. This is time-consuming

- 1024 →
- 512 →
- 256 → F', B
- 128 → E'
- 64 →
- 32 →

To make more efficient, buddy alloc can be implemented with lazy coalescing. Released block will return to its free list, but will not merge with buddy. The merge up of all avail buddies will be done when higher order block is not avail (for alloc).





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>

