

Embedded Operating System



Agenda

- **File Management**

1. Directory
2. Links
3. File System Architecture
4. File IO syscalls
5. Disk allocation & Free space management
6. Linux Ext2/3 FileSystems
7. Journaling
8. Disk scheduling algorithms

- ❖ **Reading**

1. Galvin slides (File System & IO subsystem)
2. Professional Linux Kernel Architecture (Virtual File System, Extended File System)
3. Beginning Linux Programming (File SysCalls Programming)



File Management

- **File = Data + Metadata**
 - Data --> Data blocks
 - Metadata --> Inode (FCB)
- **File System = Boot block + Super block + Inode list + Data blocks**
- **Types of Files**
 - User perspective
 - Text files
 - Archive files
 - Media files
 - Document files
 - Executable files
 - etc.



• **Kernel perspective**

- Regular files (-) (All user perspective file are regular file)
- Special files
 - Directory files (d)
 - Link files (l)
 - Pipe files (p)
 - Socket files (s)
 - Device files
 - Char device files (c)
 - Block device files (b)



• **Directory**

- From end user perspective, directory is a container which contains sub-directories and files.
- However, OS treats directory as a special file.
- The directory file contains one entry for each subdirectory or file in it.
- Each directory entry contains i-node number and name of sub-dir / file.
- terminal> ls -la -i -l /home/sunbeam.

• **Directory Listing**

- terminal> ls dirpath
- Directory access library functions (man section 3)
 - opendir()
 - readdir()
 - closedir()



- opendir()
 - Open the directory file for reading.
 - DIR *dp = opendir("dir-path");
 - arg1: dir path to be opened
 - returns: DIR pointer if dir opened successfully, otherwise NULL
- closedir()
 - Close the directory file.
 - closedir(dp);
 - arg1: DIR pointer.



• **readdir()**

- Read the next dirent from the directory file.
- `struct dirent *ent = readdir(dp);`
 - arg1: DIR pointer.
 - returns: Pointer to struct dirent, if next entry is available.
 - Returns NULL if end of dir file is reached.
- struct dirent
 - d_name --> name of file or sub-directory.
 - d_ino --> inode number of file or sub-directory



- Symbolic Link

- A symbolic link, also known as a symlink or soft link, is a special type of file that points to another file or directory.
- terminal> ln -s /path/of/target/file linkpath
- Internally use
 - symlink() syscall.
 - man symlink
 - int symlink(const char *target_path, const char *link_path)

- symlink()

- syscall A new link file is created (new inode and new data block is allocated), which contains info about the target file (absolute or relative path).
- Link count is not incremented.
- If target file is deleted, the link becomes useless.
- Can create symlinks for directories also



• Hard Link

- A hard link to a file points to the inode of the file instead of pointing to the file itself.
- This way the hard link gets all the attributes of the original file and points to the same data block as the original file.
- terminal> In targetfilepath
- Internally use link() syscall.
 - man link
 - `int link(const char *target_path, const char *link_path);`
- link() syscall
 - A new directory entry is created, which has a new name and same inode number.
 - No new file (inode and data blocks) is created.
 - Link count in the inode of the file is incremented.
 - If directory entry of target file is deleted (rm command), file can be still accessed by link directory entry.
 - Cannot create hard link for directories, because it may lead to infinite recursion (while traversing directories recursively e.g. ls -R)



- rm command

- The 'rm' means remove.
- This command is used to remove a file.
- The rm command in Linux, internally calls unlink() system call.
- `int unlink(const char *filepath);`

- unlink() syscall

- It deletes directory entry of the file.
- It decrements link count in the inode by 1.
- If link count = 0, the inode is considered to be deleted/free (updated into super-block).
- It can be reused for any new file.
- When inode is marked free, data blocks are also made free, so that they can also be reused for some new file



• Directory

- Directory permissions/mode
 - r -- can read from dir data block -- list directory contents.
 - w -- can write into dir data block -- create new files & sub-directories, remove file/sub-directory, rename file.
 - x -- enable browsing the directory -- "cd" command

SUNBEAM



- File System Architecture

- Virtual File System:

- This layer redirect file system request to the appropriate file system manager.

- File system manager:

- File system manager enables access to repective file system on the disk.
 - OS can see all partitions whose file system managers are installed in that OS.

- IO subsystem:

- Implement buffer cache and other mechanisms to speed up disk IO.

SUNBEAM



• **Windows vs Linux**

- Linux have FS mgrs for ext3/4, reiserfs, xfs, fat, ntfs, cdfs, etc.
- Hence Linux support many FS.
- Windows have FS mgrs for FAT, NTFS, CDFS.
- Hence Windows do not support Linux FS.
- However, third-party FS managers can be added into Windows to support Linux FS e.g. ext2fsd.

SUNBEAM



• File IO syscalls

- open() syscall
 - fd = open("/home/kiran/abc.txt", O_RDONLY);
 - step 1. Convert given file path into its inode number. This is called as path name translation and is done by a kernel in file from the disk into inode table in memory.
 - step 2. Inodes of all recently accessed files are kept in inode table.
 - step 3. A file position is initialized to 0 and is stored in the open file table. It also stores mode in which file is opened and pointer to the in-memory inode. Information of all files opened in the system, is maintained in this table.
 - step 4. Each process is associated with a open file descriptor table. It keeps info of all files opened by that process. This entry stores pointer to the Open FileTable entry.
 - step 5. Finally index to file desc table entry is returned, which is called as "file descriptor". All further read(), write(), lseek(), close() operations will be using this file desc.



• VFS Structures (inode table)

- struct inode unsigned long i_ino; // inode number
- loff_t i_size; // file size unsigned
- int i_nlink; // number of hard links
- umode_t i_mode; // file mode (permissions)
- atomic_t i_count; // reference count
- struct list_head i_list; // inode cache
- Device driver related
 - struct list_head i_devices;
 - dev_t i_rdev;
 - union {
 - struct pipe_inode_info *i_pipe;
 - struct block_device *i_bdev;
 - struct cdev *i_cdev;
 - };
 - struct file_operations *i_fop;



- **struct file (Open File Table)**

- unsigned int f_flags; // open() arg2
- loff_t f_pos; // current file position
- struct path f_path; // pointer to dentry
- #define f_dentry f_path.dentry
- struct list_head fu_list; // open file table
- atomic_t f_count; // reference count
- Device driver related
 - struct file_operations *f_op;



- **struct dentry** (it store in process PCB)

- struct qstr d_name; // name of file/sub-directory
- struct inode *d_inode; // pointer to the inode
- struct list_head d_lru; // dentry cache
- atomic_t d_count; // reference count

- **struct fs_struct** (it store in process PCB)

- struct dentry * root; // stores "root directory" of the process --> used for absolute path
- struct dentry * pwd; // stores "current directory" of the process --> used for relative path
- int umask; // user file mode mask -- while creating new file this mask is used.



- **struct files_struct**

- struct file * fd_array[NR_OPEN_DEFAULT];

- **struct task_struct**

- struct fs_struct *fs; // current & root directory
 - struct files_struct *files; // open file desc tables

- **Reference counting**

- Used to manage life-time of any object (in complex systems e.g. Linux kernel, ...).
 - Object has a member called as "reference count".
 - The count is incremented everytime new pointer points to the same object and decremented everytime the pointer to the object is no more used/required.
 - At any moment, reference count is number of pointers referring to the object.
 - When reference count become zero, it means no pointer is referring to the object and the object can be deleted safely



- **close() syscall**

- Decrement ref count in open file table entry (struct file).
- If ref count drops to zero, OFT entry is deleted (from OFT).

- **read() syscall**

- `count = read(fd, buf, length);` -- syscall api `sys_`
- `read(fd, buffer, length)` -- syscall implementation
 - `vfs_read(file, buffer, length, inode)` -- Virtual file system
 - Logical FS considers file as sequential set of bytes and set of blocks.
 - Example: block size = 4096 and file size = 20000 bytes, then number of blocks = 5 (0 to 4).
 - If current file position = 10000, then reading file block = 2
 - `ext3_read(file, inode, file_block)` -- File system manager
 - Refers inode and file disk block corresponding to the file block.
 - check buffer cache
 - -- if disk block found. if found, return it;
 - otherwise call disk driver to read that disk block from
 - `disk_read(disk_device, disk_block)`
 - -- device driver Read appropriate sectors and made it available into buffer cache.
 - The current process is blocked/sleep while disk read operation is in progress.





Thank you!

Kiran Jaybhavne

email – kiran.jaybhavne@sunbeaminfo.com

