

(Non-Preemptive)

Process	Arrival	CPU Burst
P1	0	7
P2	2	4
P3	4	1
P4	5	4

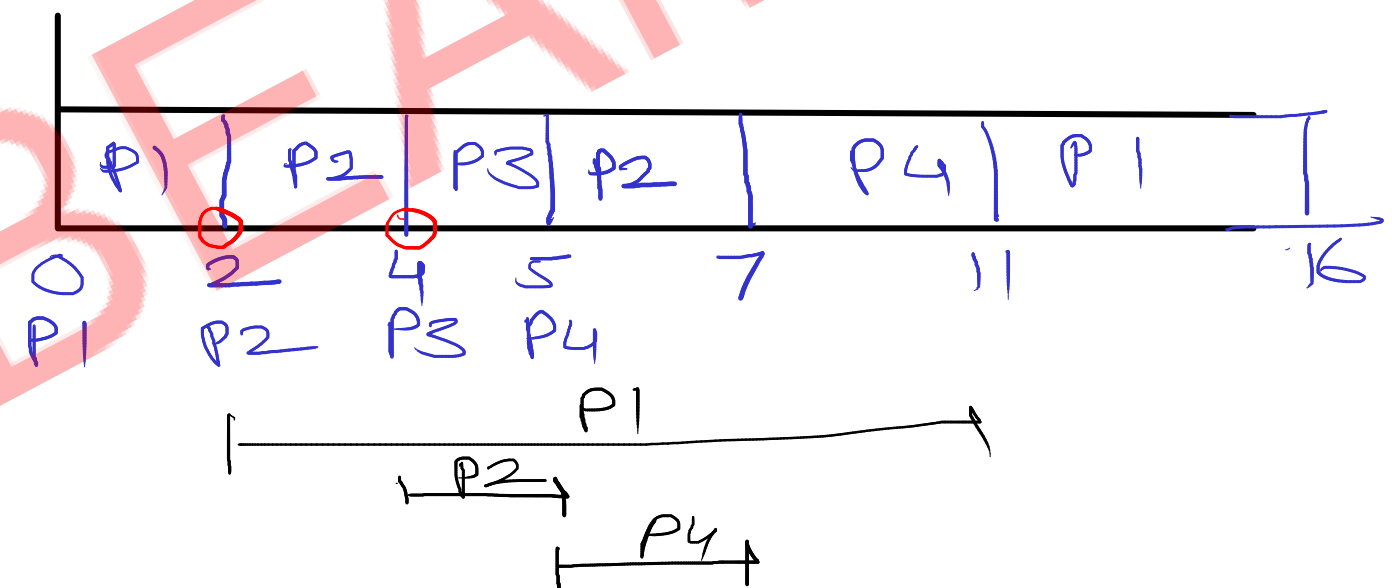
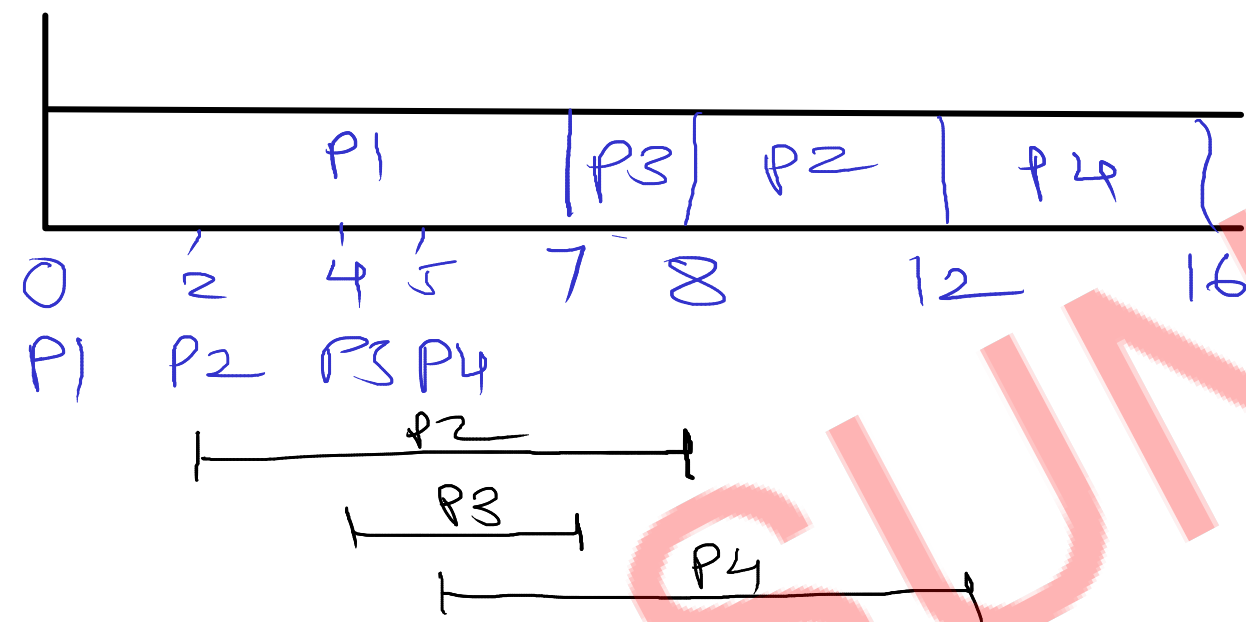
WT RT TAT
 0 0 7
 6 0 10
 3 0 4
 7 7 11

SJF (Shortest Job First)
 (Preemptive)
 (Shortest Remaining Time First)

Process	Arrival	CPU Burst
P1	0	7
P2	2	4
P3	4	1
P4	5	4

Remain
time

WT RT TAT
 9 0 16
 1 0 5
 0 0 1
 2 2 6



Starvation

- due to longer CPU burst, process will not get enough CPU time to execute
- there is no solution for starvation

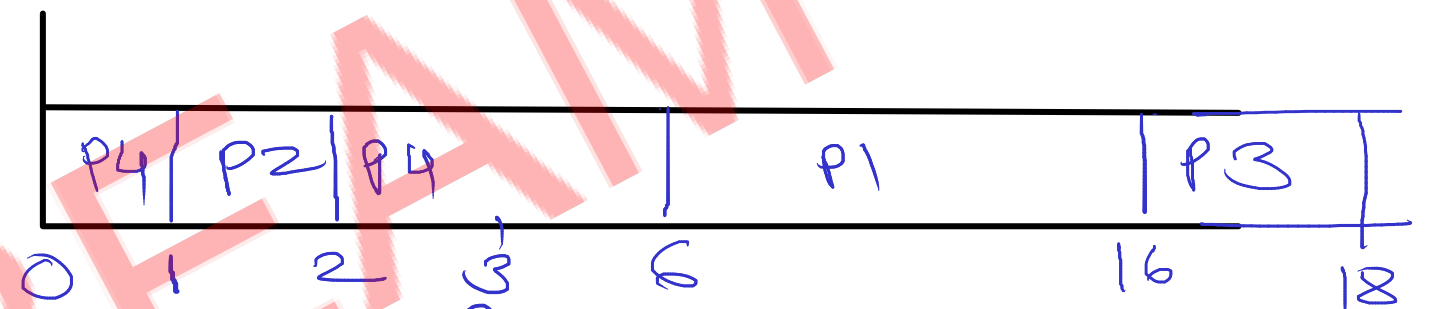
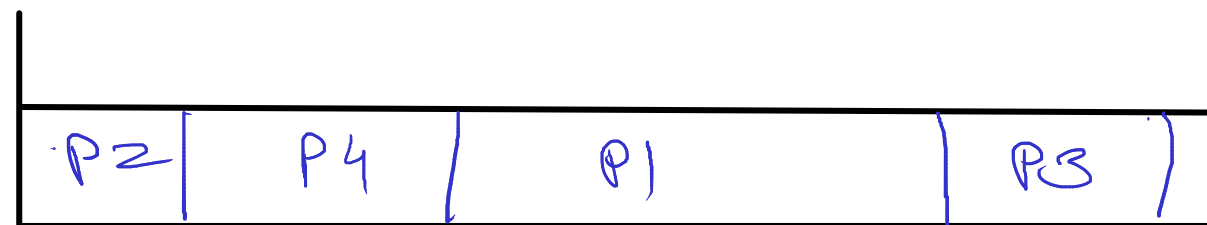
Priority

(Non-Preemptive)

Process	Arrival	CPU Burst	Priority	WT	RT	TAT
P1	0	10	3	6	6	16
P2	0	1	1 (H)	0	0	1
P3	0	2	4 (L)	16	16	18
P4	0	5	2	1	1	6

(Preemptive)

Process	Arrival	CPU Burst	Priority	WT	RT	TAT
P1	0	10	3	6	6	16
P2	1	1	1	0	0	1
P3	3	2	4	13	13	15
P4	0	5	2	4	0	6



0
P1
P2
P3
P4

0
P1
P2
P3
P4

P1 (6)
 P2 (9)
 P3 (7)
 P4 (5)
 P5 (8)
 P6 (7)
 P7 (6)

P1
 P4
 P3
 P7
 P6
 P5
 P2

Starvation

- due to less priority of process, it is not getting enough time to execute on CPU

Aging

- increase the priority of process (whose priority is less) gradually

(pure preemptive)

RR (Round Robin)

- CPU time is divided into time slices (Time Quantum)
- for every time quantum one process is scheduled

Time quantum = 20

Process	CPU Burst
P1	53
P2	17
P3	68
P4	24

Remain time

33, 13, 0X

0X

48, 28, 8X

4, 0X

waiting time

0 + 57 + 24

20

37 + 40 + 17

57 + 40

Response time

0

20

37

57

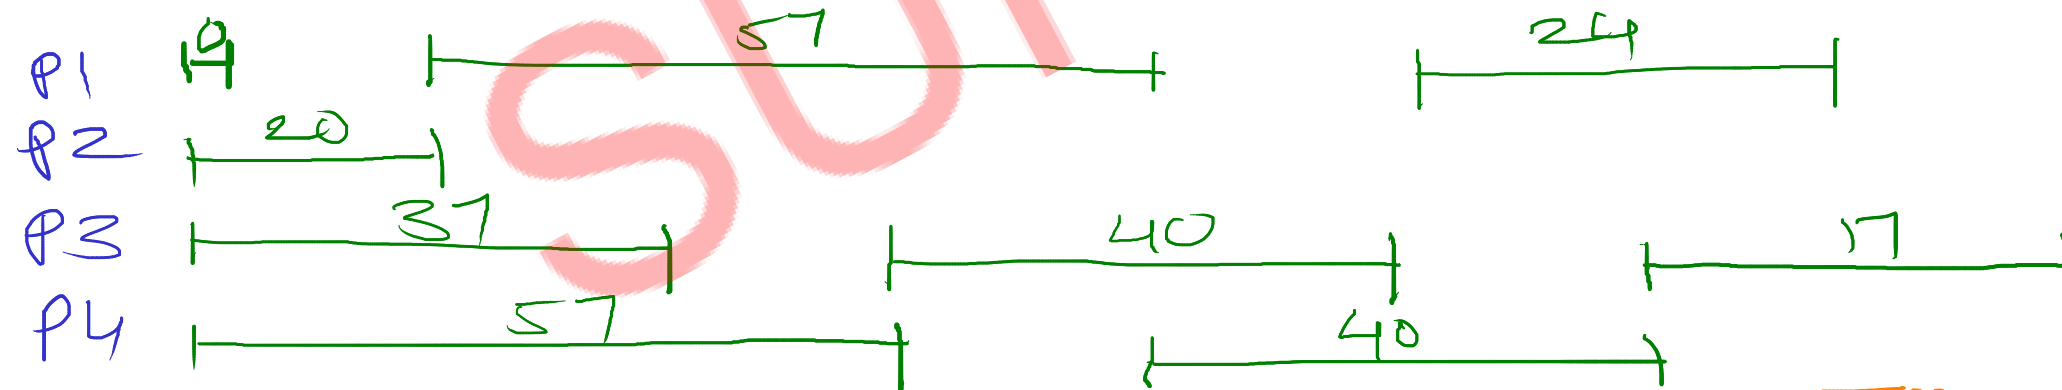
TAT

134 - 0 = 134

37 - 0 = 37

162 - 0 = 162

121 - 0 = 121



first waiting time is always response time

TG = 100

↳ behave like FCFS

TG = 4

↳ CPU overhead will increase

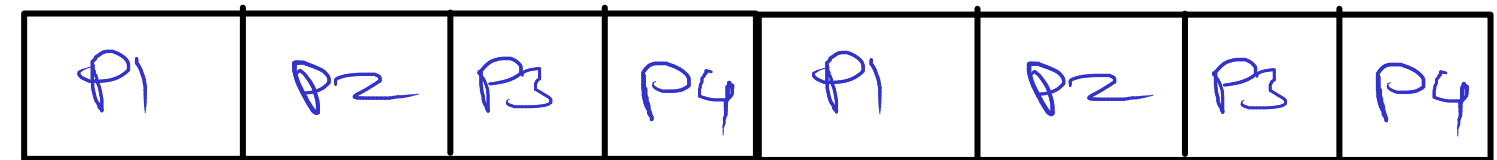
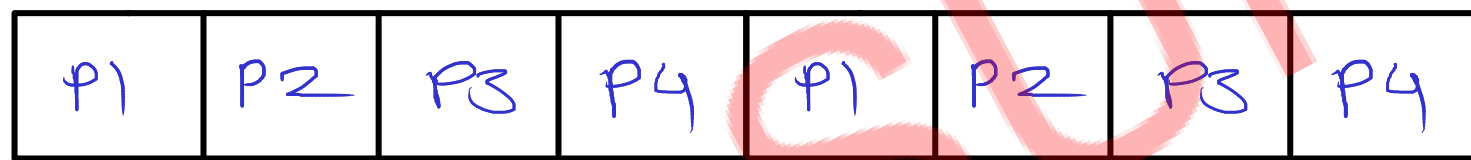
Fair Share

- CPU time is divided into time slices (epoch)
- some share of each epoch is given to the processes which are in ready queue.
- share is given to the process on the basis of their priority
- priority of every process is decided by its nice value
- nice values range ---> -20 to +19 (40 values)
 - * -20 - highest priority
 - * +19 - lowest priority

Process	Nice Value
P1	10
P2	10
P3	10
P4	10

Epoch - 100

Process	Nice Value
P1	5
P2	5
P3	10
P4	10



1. `ps -e -o pid,ppid,ni,cmd`

2. `nice -n 10 ./demo01.out`

3. `sudo nice -n -10 ./demo01.out`

4. `renice -n 10 -p <pid>`

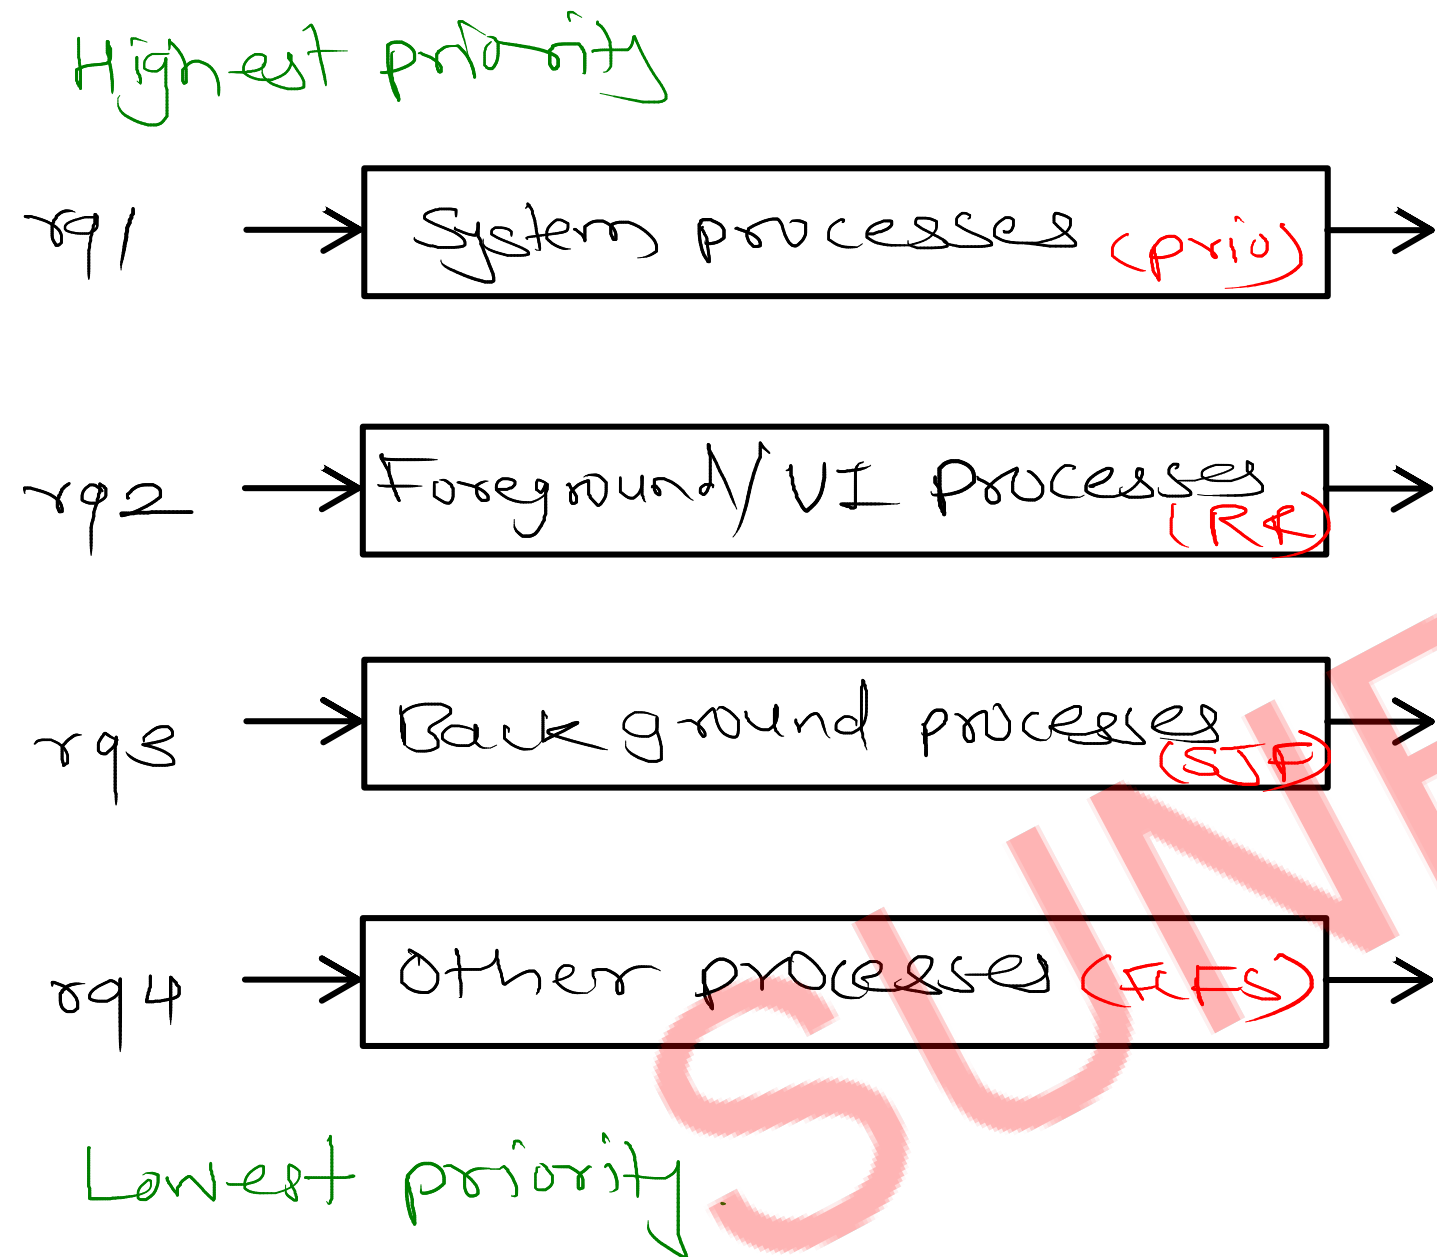
- to see the nice value

- to assign nice value to the process (decrease prio)

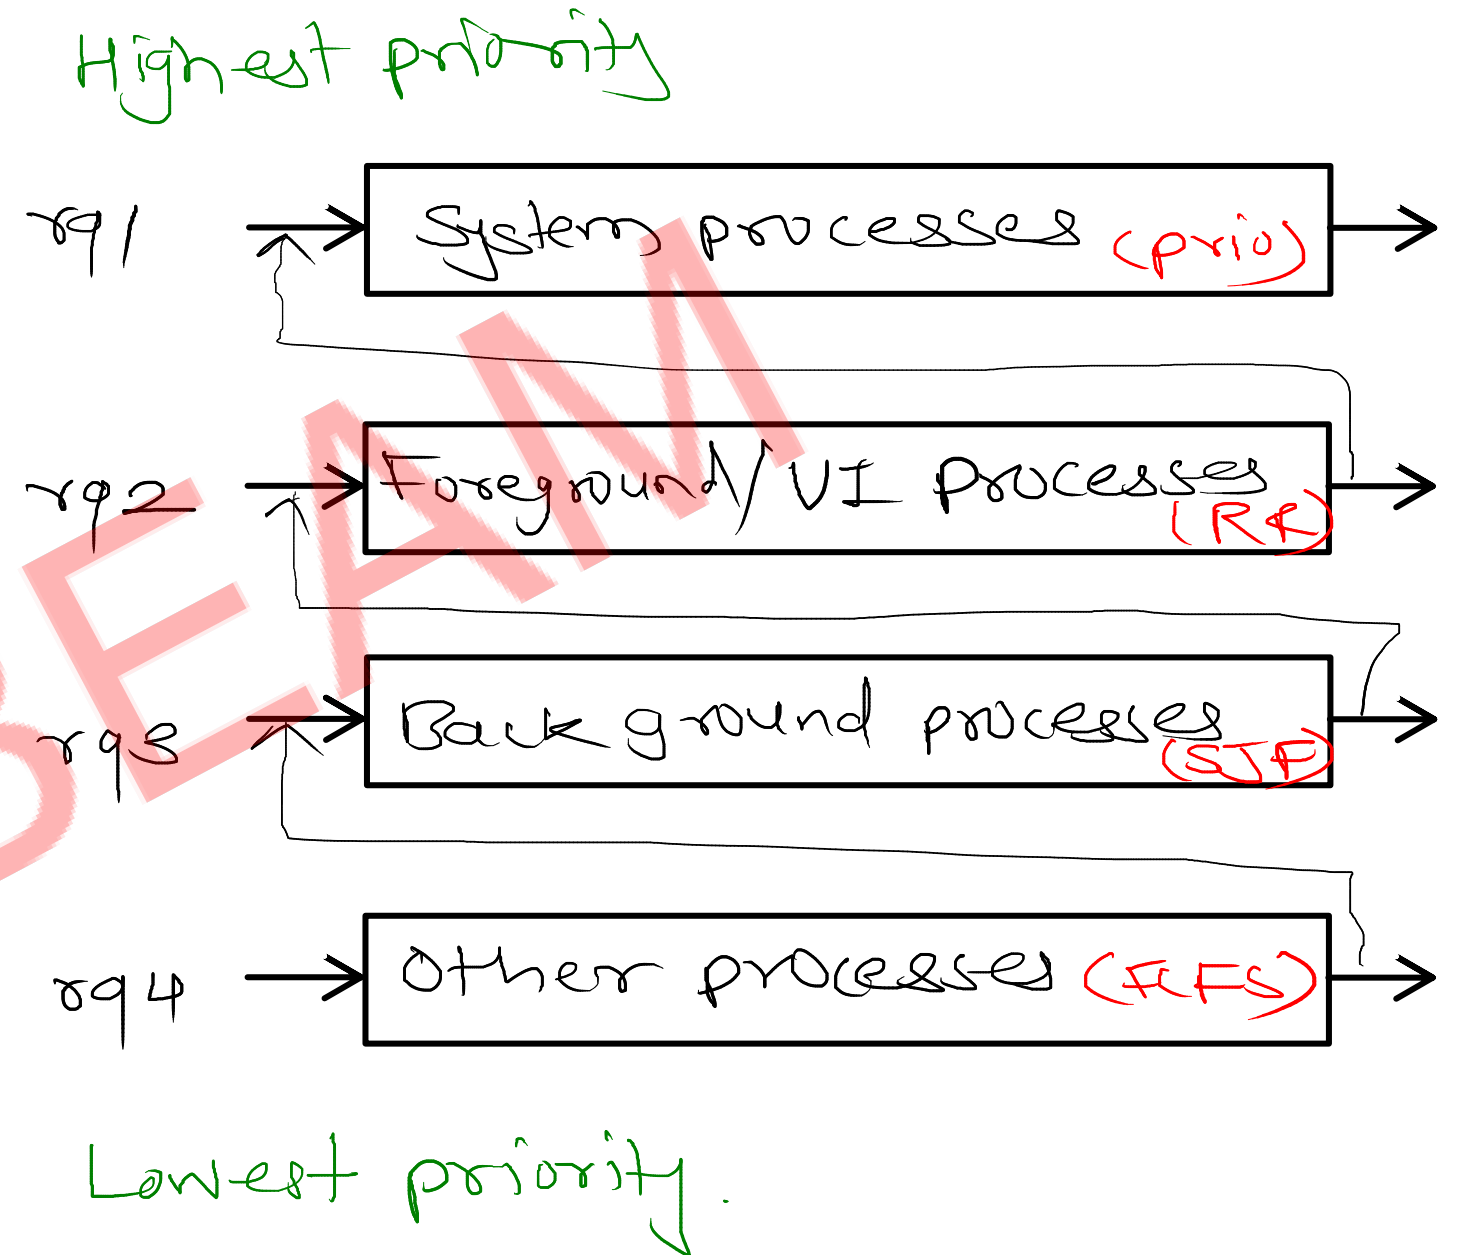
- to assign nice value to the process (increase prio)

- to change nice value of process at runtime

Multi level Ready Queue



Multi level Feedback Ready Queue



— fixed priority preemptive scheduling

Latencies

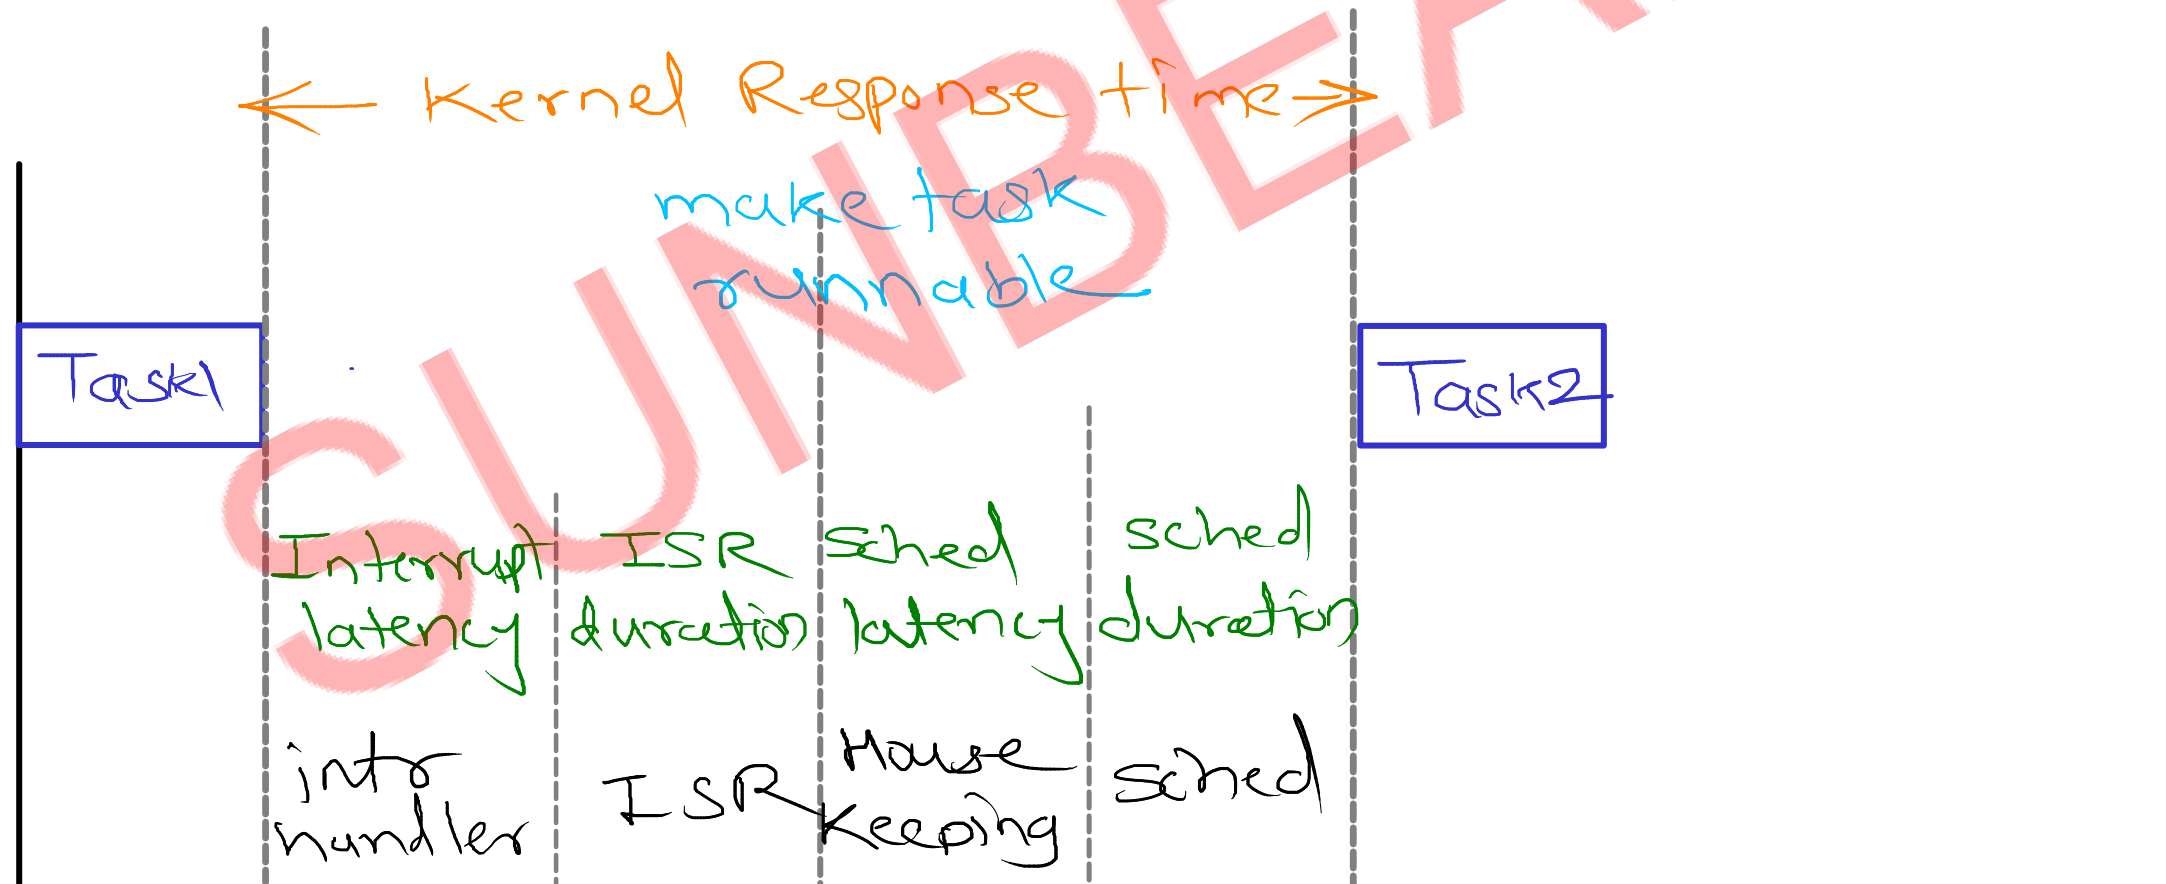
Interrupt latency - time from arrival of interrupt to start of ISR execution

Scheduler latency - time from completion of ISR to start of scheduler

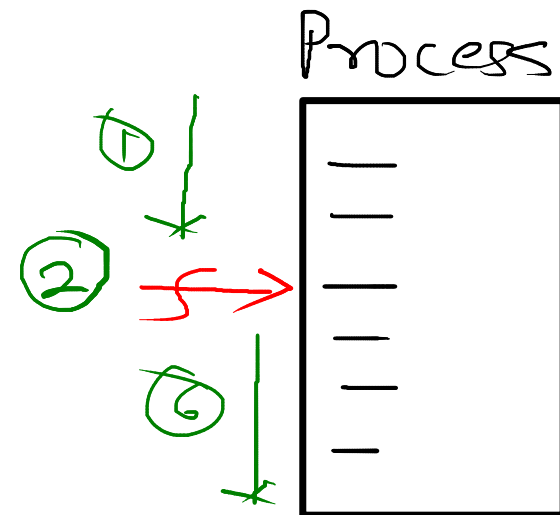
Kernel Response Time (Scheduling Latency)

- time to start next process / task

= Interrupt latency + ISR duration + Scheduler latency + Scheduler duration



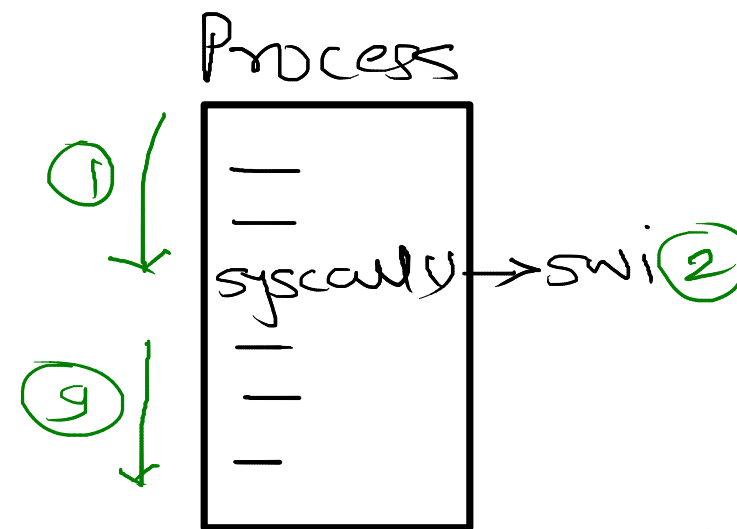
User Space Preemption



interrupt_handler()
{

- 1) save exe context of current process
 - 2) Find address of ISR from IVT
 - 3) call ISR
 - 4) pid = scheduler()
 - 5) dispatcher(pid)
- }
-
- The diagram shows a flow of control. A green arrow labeled '3' points down from the 'call ISR' step to a blue line. From this blue line, a green arrow labeled '4' points down to the 'pid = scheduler()' step. A green arrow labeled '5' points down from the 'dispatcher(pid)' step to the closing brace of the interrupt handler. A green arrow labeled 'X' points from the 'pid = scheduler()' step to the 'dispatcher(pid)' step. A blue line also connects the 'call ISR' step to the 'pid = scheduler()' step.

Kernel Space Preemption



interrupt_handler()

{

1) save exe context of current process

2) Find address of ISR from IVT

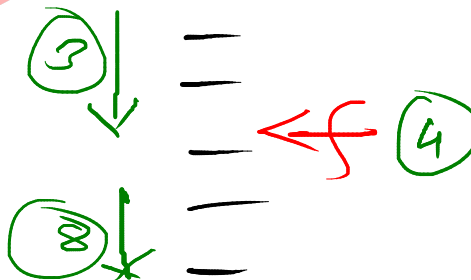
3) call ISR

4) pid = scheduler()

5) dispatcher(pid)

}

syscall()

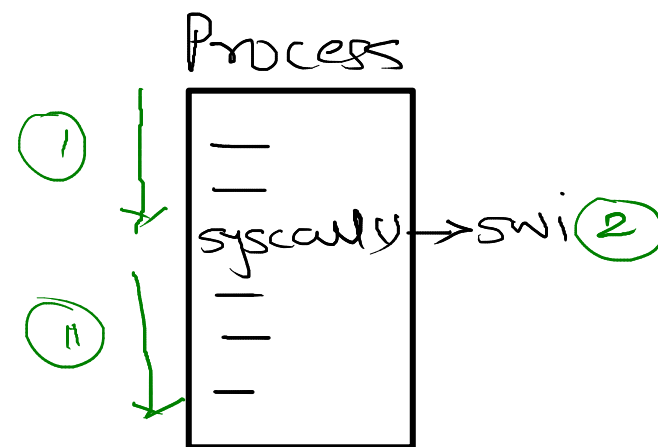


ISR()

⑥ ↓

ⓧ

User Space Preemption



swi_handler()

```

{
  - get system call number;
  3 - find actual system implementation
    disable_intr()
    - syscall()
    7 enable_intr()
    - pid = scheduler();
    - dispatcher();
}

```

interrupt_handler()

```

{
  1) save exe context of current process
  8 2) Find address of ISR from IVT
    3) call ISR
    4) pid = scheduler()
    5) dispatcher(pid)
}

```

syscall()

```

{
  4 -
  5 -
  6 -
  7 -
}

```

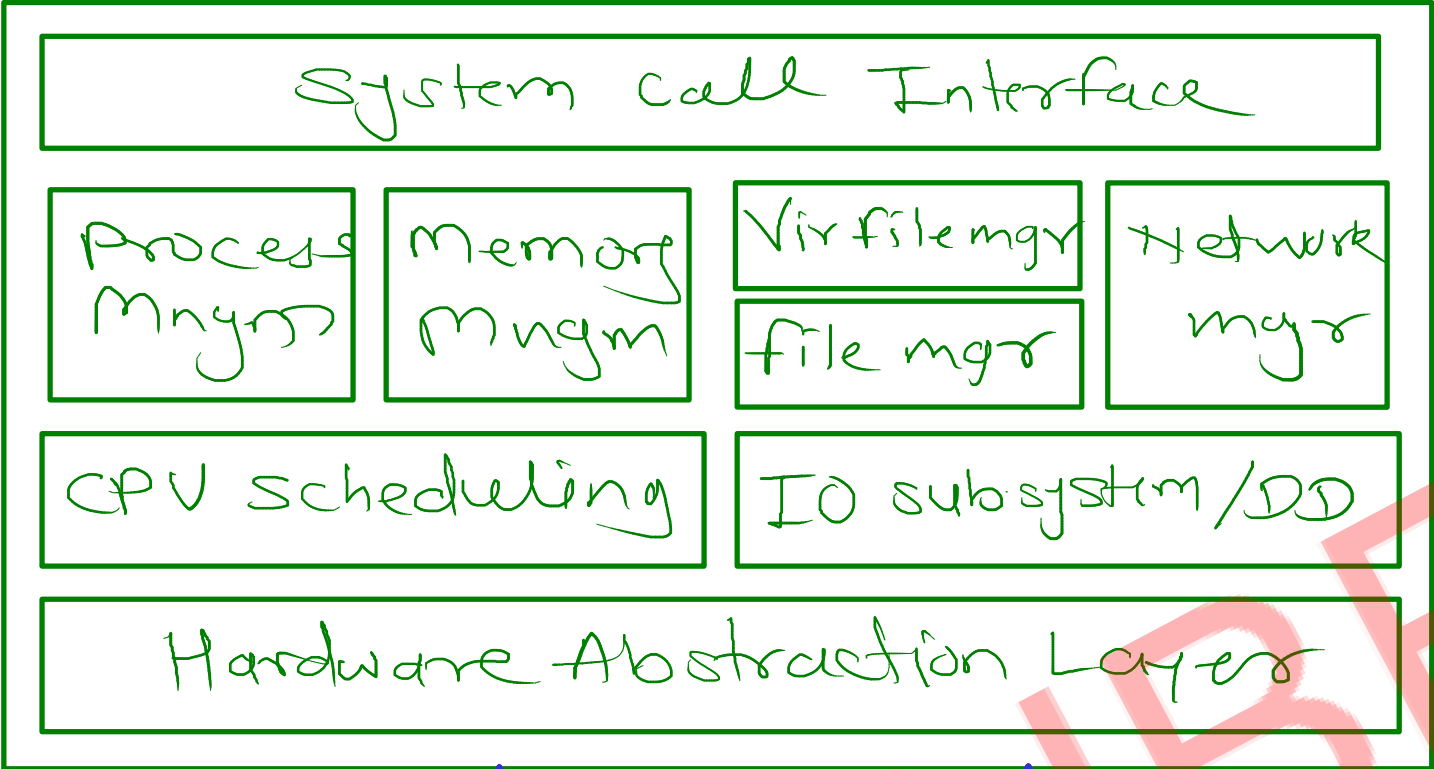
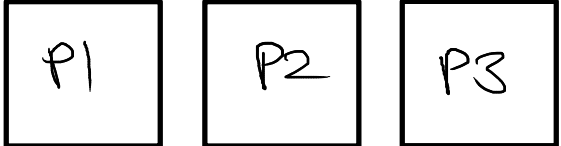
Interrupt is masked

ISR()

9

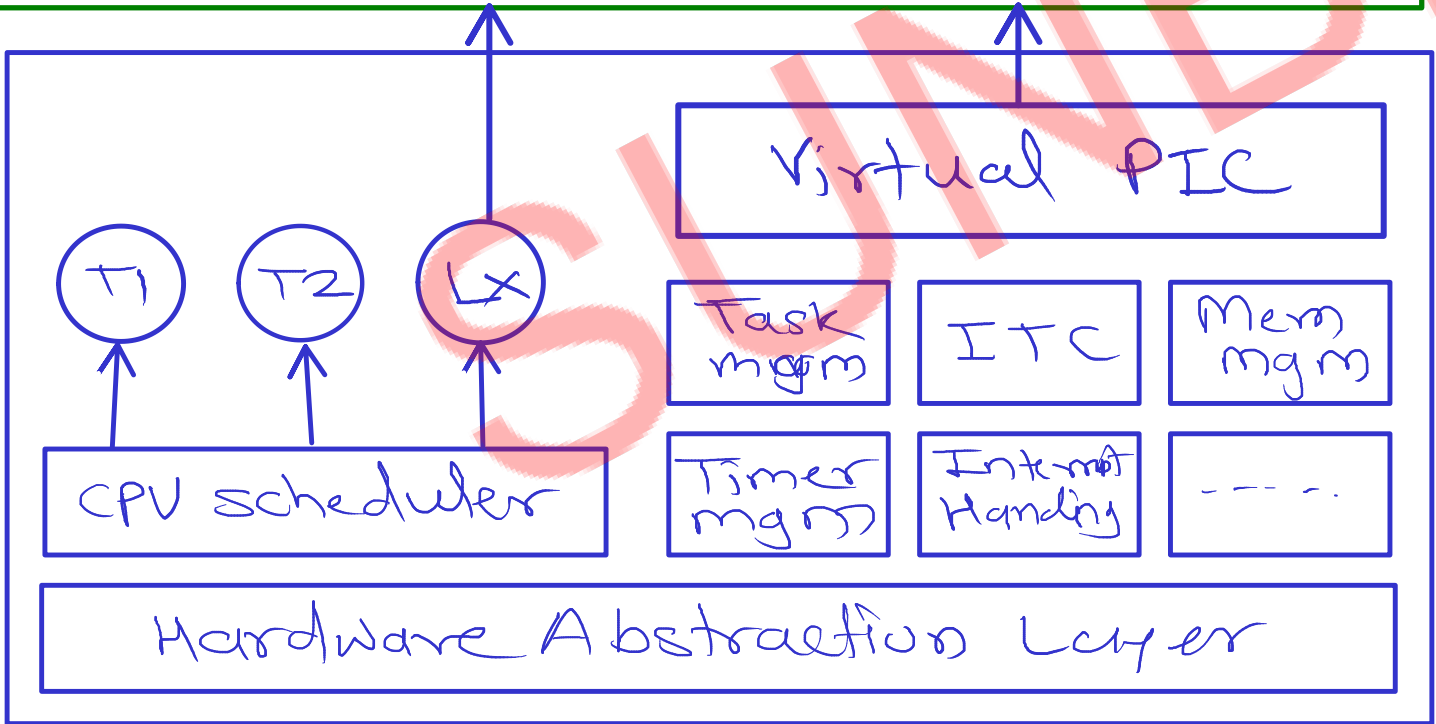
X

Dual Kernel Approach (Linux Based RTOS)

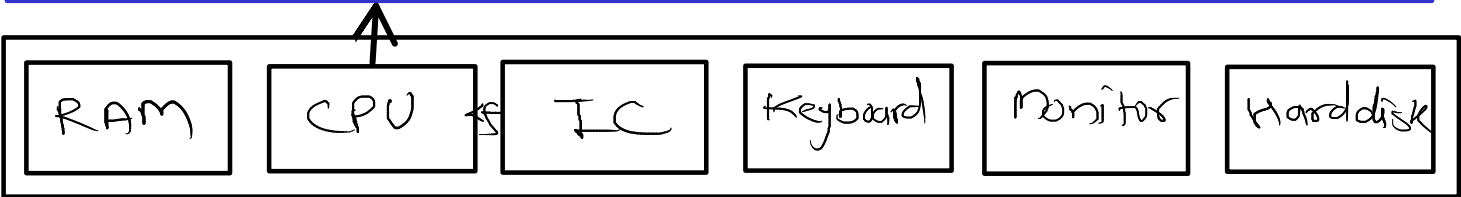


Linux
Kernel

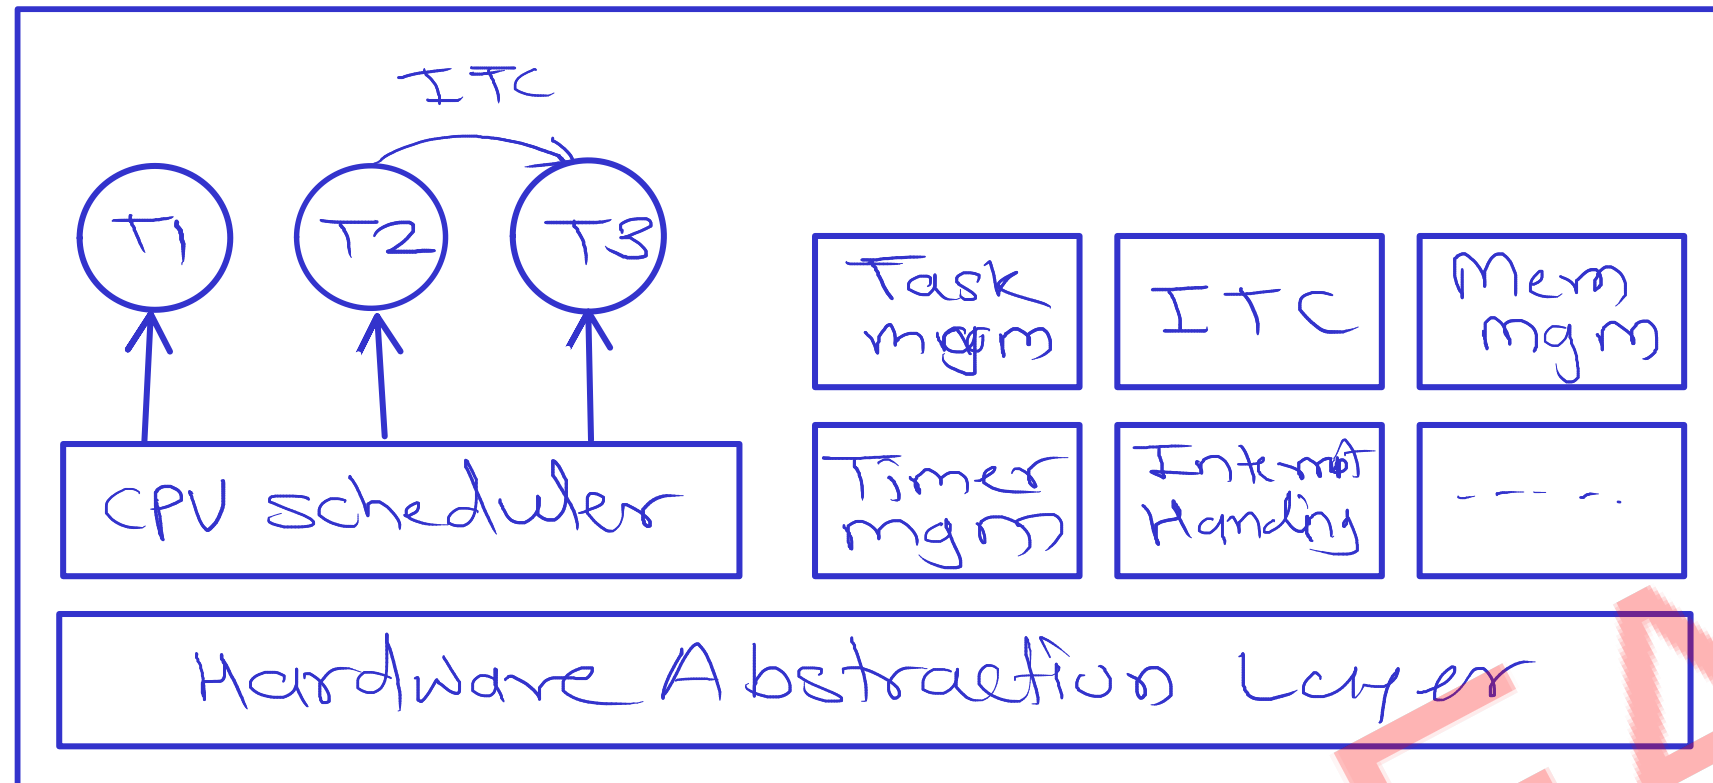
Xenomai
RTAI



RTOS
Kernel



Embedded RTOS



Embedded
RTOS
Kernel

FreeRTOS
VxWorks
QNX
uCOS

