

Linux

Linux

- ❑ Used on desktops, servers, smartphones, embedded devices
- ❑ As of Nov 2014, 97% of all HPC system ran Linux

Scheduling through Version 2.5

- ❑ Prior to kernel version 2.5, Linux ran a variation of standard UNIX scheduling algorithm
- ❑ Version 2.5 moved to constant order $O(1)$
- ❑ Linux uses the term "task" for process

Linux 2.5 Scheduling

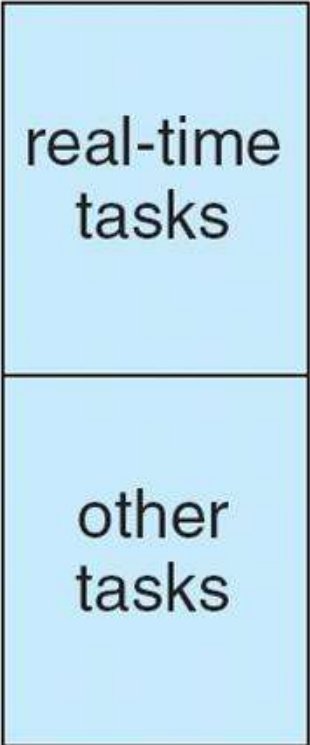
□ Goals

- Implement scheduling algorithms in $O(1)$ time.
- Optimize for the common case of only one or two runnable processes, yet scale well to multiple processors, each with many processes.

□ Linux uses the term **task** to refer to a process/thread

□ Linux priority levels are on the next slide

Linux 2.5 - Priority Levels

<u>numeric priority</u>	<u>relative priority</u>		<u>time quantum</u>
0	highest		200 ms
• • • 99			
100			
• • •			
140	lowest		10 ms

Priorities and Time-slice length

Linux 2.5 - Priorities

- ❑ Priorities 0-99 for real-time processes
- ❑ Priorities 100-139 for normal (user) processes

Linux 2.5 Scheduling

- ❑ A process is considered for execution on the CPU as long as it has time remaining in its time slice (quantum)
- ❑ When a task has exhausted its time slice, it is considered **expired** and is not eligible for execution again until all other tasks have all exhausted their time slice

Linux 2.5 Data Structures

- ❑ The kernel maintains a list of all runnable tasks in a **runqueue** data structure
- ❑ A runqueue consists of two priority arrays:
 - **Active**: Contains all tasks with time remaining in their time slices
 - **Expired**: Contains all expired tasks
- ❑ The active, expired queues are indexed according to priority

Linux 2.5 Data Structures



List of Tasks Indexed According to Priorities

Linux 2.5 Scheduling

- ❑ How is a process selected for execution?
 - The scheduler finds the highest-priority queue with a runnable process
 - Finds the first process on the queue

- ❑ What happens to a running process that does not complete its time quantum?
 - When that process returns to the **ready** state it goes into the active array

Linux 2.5 Scheduling

- When there are no more processes in the **active** array, the **expired** array becomes the **active** array and the **active** array becomes the **expired** array.
- A process's priority (queue) is recalculated when the process has exhausted its time quantum and is to be moved to the expired array
 - Calculation considers time spent waiting for I/O